

COE3DQ5 Project Description 2024

Hardware Implementation of an Image Decompressor

1 Introduction

1.1 Objective

In this project, students will gain experience in digital system design by implementing the custom McMaster Image Compression revision 18 (.mic18) image compression specification in hardware. Compressed data for a 320 x 240 pixel image will be delivered to the Altera DE2-115 board via the universal asynchronous receiver/transmitter (UART) interface from a personal computer (PC) and stored in the external static random access memory (SRAM). The image decoding circuitry will read the compressed data, recover the image using a custom digital circuit you designed, and store it back to the SRAM, from which the video graphics array (VGA) controller will read it and display it to the monitor.

1.2 Preparation

- Revise the first five labs, especially finite state machines, VGA, SRAM and UART interfaces
- Read this document and get familiarized with the C source code

1.3 Organization of this Document

Section 2 discusses how an image is compressed with the .mic18 specification. Section 3 demonstrates image decompression by example with a 16x16 image. Section 4 presents the project guidelines. The project is divided into three milestones.

2 Image Compression

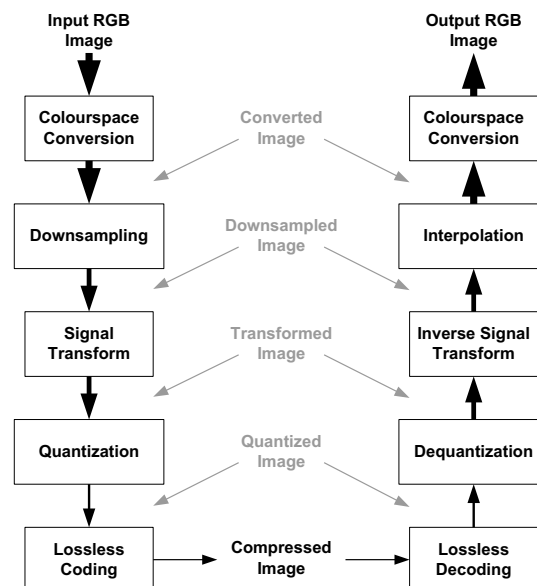


Figure 1. Flow of basic image compression (down the left) and decompression (up the right).

Figure 1 shows the conceptual flow for basic image compression and decompression. The size of the arrow roughly indicates the relative amount of data at each stage. In general, for image compression, we start with an image in the standard red/green/blue (RGB) format. This is converted to YUV format, where Y represents brightness, and U and V are chrominance components (which contain colour information). The YUV image is downsampled and then transformed into the frequency domain. This result is quantized, and finally, we apply lossless encoding to obtain the compressed information bitstream. Decoding proceeds in the reverse order: lossless decoding and dequantization (restores the frequency domain image), then the inverse signal transform (which brings the image back to the spatial domain), and finally, upsampling and colourspace conversion (produces an RGB image).

Compression enables smaller file sizes, which reduces storage and transmission costs. We convert the RGB image to YUV because the human eye is more sensitive to brightness than colour, so if we want to obtain more lossy compression (by reducing the quality), it will be more effective to compress the colour information. In the RGB format, there is no immediate way to compress colour more than brightness. YUV information is compressed by downsampling U and V in the .mic18 specification. Additionally, the human eye is less sensitive to high-frequency components in an image (compared to low frequencies). Transforming into the frequency domain enables us to compress higher-frequency components more than lower-frequency ones. We compress frequency information by applying quantization. Finally, we apply lossless compression as a final effort to reduce the file size.

Downsampling and quantization are not identically reversible, so some information is lost during compression. The decompressed image only approximates the original image. A good compression scheme leverages this loss of information to reduce the compressed file size but will attempt to allow information loss only in a way that does not drastically affect the quality of the decompressed image.

The focus of this project is the hardware implementation of a compression scheme. The .mic18 specification is not meant to rival existing compression standards for images. Simplifying choices have been made in developing the specification to facilitate a straightforward hardware implementation. This limits the capabilities in terms of the achievable compression/quality tradeoff. Even so, the concepts described above still apply in a general sense to the existing compression schemes. The following subsections will discuss each stage of image compression in more detail.

2.1 Colourspace Conversion and Downsampling




Figure 2. Colourspace conversion (RGB to YUV) and horizontal downsampling of U and V.

Figure 2 shows how a 320x240 image can be decomposed into its red, green, and blue components. Each pixel (composed of 1 red, 1 green, and 1 blue value) is independently transformed into the YUV colourspace. Each resulting pixel contains 1 Y (brightness), 1 U (blue chroma), and 1 V (red chroma). The chroma components contain only colour information and no brightness. In Figure 2, the bottom of the motorcycle is red, and the corresponding bright area in the V image indicates there are large values at these pixels, representing a strong red colouring. Note this says nothing about whether it is bright red or dark red, as the brightness comes only from Y.

Each RGB pixel is converted to a YUV pixel by performing the following matrix multiplication:

$$\begin{bmatrix} Y \\ U \\ V \end{bmatrix} = \begin{bmatrix} 0.257 & 0.504 & 0.098 \\ -0.148 & -0.291 & 0.439 \\ 0.439 & -0.368 & -0.071 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix} + \begin{bmatrix} 16 \\ 128 \\ 128 \end{bmatrix} \quad (1)$$

The terminology YUV originates from the video transmission field, and different notations (depending on the video/colour/transmission standard) can be found in the literature (such as YIQ, YPbPr, YCbCr).

The human eye is much more sensitive to brightness than colour (due to more rods than cones in the eye). For this reason, we choose to downsample U and V (the chroma components) to reduce the amount of data while having little effect on how the image is perceived. We horizontally downsample by 2, so the U and V images become half the original width (no vertical downsampling so the height is unchanged). Therefore, only 2/3 of the total data remains.

2.2 Signal Transform (into the Frequency Domain) and Quantization

The next compression stage involves transforming the image from the spatial domain to the frequency domain. High spatial frequencies (brightness/colour transitions that occur in a small amount of space) are not as noticeable to the human eye, so they may be reduced in accuracy (or even removed) without drastically affecting the overall quality of the image.

In the .mic18 specification, we use a 2-dimensional discrete cosine transform (DCT), and we work on blocks of 8x8 pixels (each 8x8 block is processed independently; we do not consider images with heights and widths not divisible by 8). The 2D DCT is computed with the following matrix multiplications:

$$\mathbf{S}' = \mathbf{CSC}^T \text{ where each element } [\mathbf{C}]_{i,j} = \begin{cases} \sqrt{\frac{1}{8}} \cos\left(\frac{i\pi}{8}(j+0.5)\right) & \text{if } i = 0 \\ \sqrt{\frac{2}{8}} \cos\left(\frac{i\pi}{8}(j+0.5)\right) & \text{otherwise} \end{cases} \quad (2)$$

\mathbf{S} is the original 8x8 block (spatial domain), and \mathbf{S}' is the transformed 8x8 block (frequency domain). In Equation 2 above, i and j are the row and column indexes, respectively, each goes from 0 to 7 inclusive.

To obtain compression, we quantize \mathbf{S}' . Each element in \mathbf{S}' is divided by a (possibly) different quantization coefficient and then rounded to the nearest integer. A larger coefficient means that after division, we obtain a smaller magnitude integer, which requires fewer bits to be stored. The top left corner corresponds to frequency 0. The frequencies increase as we move down and to the right in the matrix. Thus, we should choose smaller quantization coefficients for the upper left part of the matrix and large coefficients for the lower right part to exploit the perception of the human eye.

For typical images, many values of S' will become 0 after quantization, and these zeros can be stored efficiently, as will be shown in Section 2.3. Even if nonzero, the division results in an integer that requires fewer bits to be represented. Two possible quantization matrices can be used in .mic18:

$$Q_0 = \begin{bmatrix} 8 & 4 & 8 & 8 & 16 & 16 & 32 & 32 \\ 4 & 8 & 8 & 16 & 16 & 32 & 32 & 64 \\ 8 & 8 & 16 & 16 & 32 & 32 & 64 & 64 \\ 8 & 16 & 16 & 32 & 32 & 64 & 64 & 64 \\ 16 & 16 & 32 & 32 & 64 & 64 & 64 & 64 \\ 16 & 32 & 32 & 64 & 64 & 64 & 64 & 64 \\ 32 & 32 & 64 & 64 & 64 & 64 & 64 & 64 \\ 32 & 64 & 64 & 64 & 64 & 64 & 64 & 64 \end{bmatrix} \quad Q_1 = \begin{bmatrix} 8 & 2 & 2 & 2 & 4 & 4 & 8 & 8 \\ 2 & 2 & 2 & 4 & 4 & 8 & 8 & 16 \\ 2 & 2 & 4 & 4 & 8 & 8 & 16 & 16 \\ 2 & 4 & 4 & 8 & 8 & 16 & 16 & 16 \\ 4 & 4 & 8 & 8 & 16 & 16 & 16 & 32 \\ 4 & 8 & 8 & 16 & 16 & 16 & 32 & 32 \\ 8 & 8 & 16 & 16 & 16 & 32 & 32 & 32 \\ 8 & 16 & 16 & 16 & 32 & 32 & 32 & 32 \end{bmatrix} \quad (3)$$

S' is divided element-wise by one of the quantization matrices (suppose we pick Q_0) and then it is rounded to the nearest integer:

$$Z = \text{round}(S' ./ Q_0) \quad (4)$$

If we look at each element of Z :

$$[Z]_{i,j} = \text{round}([S']_{i,j} / [Q_0]_{i,j}) \quad (5)$$

Z will later be losslessly coded.

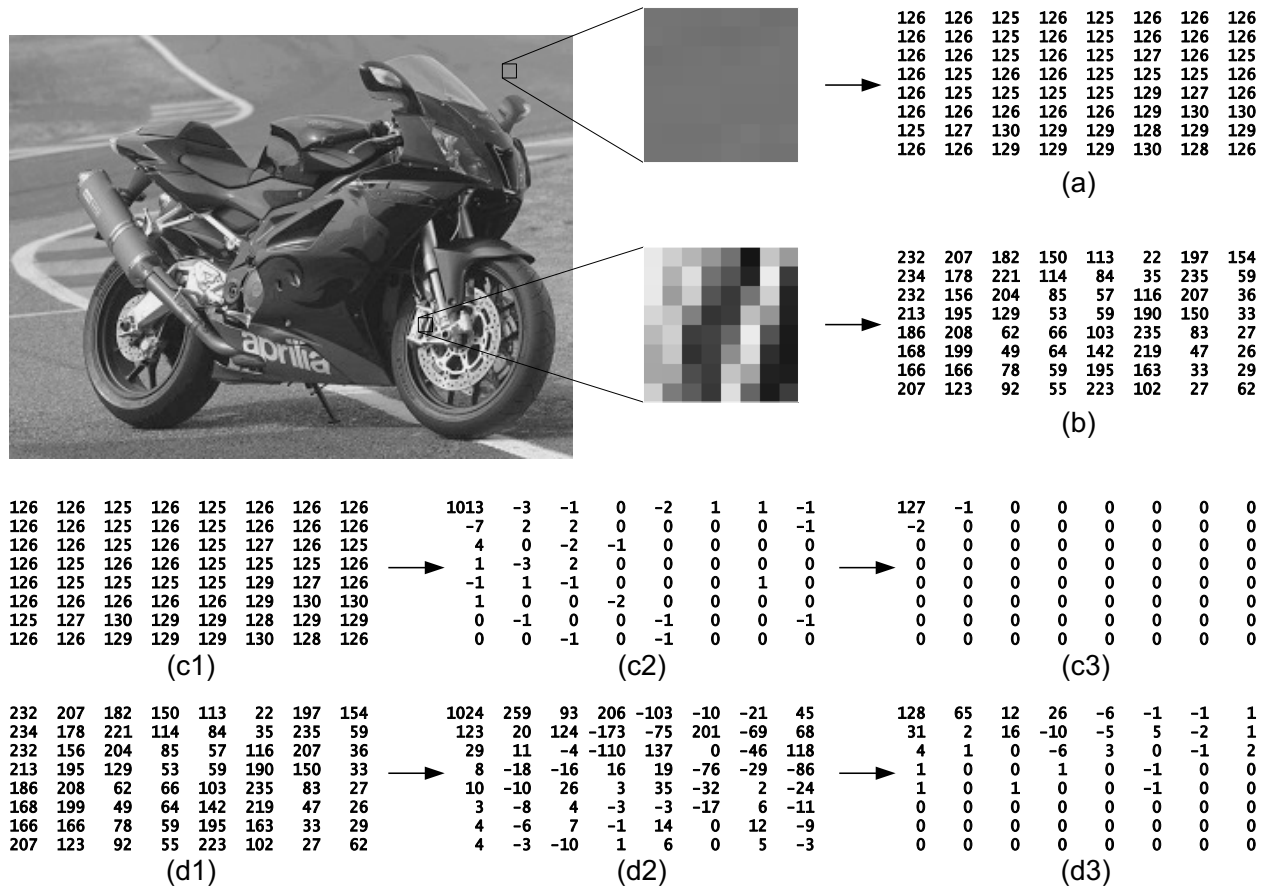


Figure 3. DCT (c1→c2 and d1→d2) and quantization with Q_0 (c2→c3 and d2→d3) of 2 blocks of Y data.

After dividing each colour plane (Y, U, and V) into blocks of 8x8 pixels, we perform DCT and quantization to each block. Figure 3 shows this on two blocks of Y data. The image is 320x240, so there will be $(320/8)*(240/8) = 1200$ Y blocks and 600 blocks each of downsampled U and downsampled V (since these planes are half as wide). All the Y blocks are processed, followed by all the U and finally all the V blocks. Within each colour plane, we go across a row of blocks first, then down to the next row. Note that one row of blocks is 8 pixels high.

The block in Figure 3(a) is nearly uniform, so after DCT, in Figure 3(c2) there is a large value for frequency 0 and small values for the high frequencies. After quantization, we end up with mostly zeros, which can be stored efficiently. Many images contain nearly uniform blocks, so the DCT followed by quantization provides a way to represent the image with much less data. Although this leads to reduced quality of the decompressed image, the major features of the original block are still preserved.

The block in Figure 3(b) contains more variation due to the details of the motorcycle's wheel. The DCT and quantization results are shown in Figure 3(d2) and Figure 3(d3) respectively. In the presence of more details (captured by the higher frequencies), more data is needed to represent the block (compared to the block from Figure 3(a)). Even so, the amount of data (in Figure 3(d3)) is reduced from that required to represent the intensity of each pixel explicitly (the block in Figure 3(b)).

2.3 Lossless Coding

In the final stage of decoding, we need a mechanism for efficiently storing the data after DCT and quantization. We start with a simple 8-byte header, which indicates the image's width and height and which quantization matrix was used. After this, the losslessly coded blocks are packed into the bitstream in the same order described at the top of this page.

To losslessly code one block, the elements of Z (see Equation 4) are read out of the block in the order shown in Figure 4. This is known as the scan pattern.




Figure 4. Scan pattern (order of values to be read from an 8x8 block).

Many images naturally have small values in the high frequencies, so after computing the DCT and quantizing, we typically have many zeros in the lower right part of the matrix. The above scan pattern, standard for many mainstream video/image compression schemes, deals with the low frequencies first (upper left) and the high frequencies last (lower right). This is beneficial because, in this order, we can get many consecutive zeros (from the high-frequency part), which can be efficiently coded. For each 8x8 block, as we follow the order of the scan pattern, values are encoded as follows, and a variable length prefix-coded header indicates what data will follow (a higher priority is given to the cases higher in the list):

1. Only zeros remain in the block, output the ZEROS_TO_END header (bits 1010).
2. A run of zeros exists but some nonzero coefficients still remain in the block. Output the ZERO_RUN header (bits 11) and then if
 - a. The run of zeros is 8 or more in length, output the 3 bits 000, now the run of zeros will be 8 elements shorter, so repeat step 2 if necessary.
 - b. The run of zeros is less than 8 in length (1-7 inclusive), output a 3 bit unsigned number indicating the length of the run.
3. There are two consecutive values between -4 and 3 inclusive, output the DOUBLE_3_BIT header (bits 00) followed by the first value (3 bits signed) and then the second value (another 3 bits signed).
4. There is a value between -4 and 3 inclusive, output the SINGLE_3_BIT header (bits 01) followed by the 3 bit signed value.
5. There is a value between -32 and 31 inclusive, output the 6_BIT header (bits 100) followed by the 6 bit signed value.
6. There is a value between -256 and 255 inclusive, output the 9_BIT header (bits 1011) followed by the 9 bit signed value.

In this entire document, it is implied that the signed representation is the 2's complement. Likewise unsigned is just the usual binary format. In the encoding, note that if a 3 bit value is followed by a run of 2 or more zeros then the 3 bit value is coded alone, but if it is followed by a single zero then the 2 values are coded together (the 3 bit value followed by the 3 bits 000).

Figure 5 shows the coding of the final (quantized) blocks from Figure 3, in both cases the blocks are scanned according to Figure 4. For the first block, the first value (127) needs to be coded as 9 bits signed (case 6 above), so we output 1011 (the header) followed by 001111111 (the 9 bit signed representation of 127). The next 2 values are -1 and -2 which can be coded together as 3 bit signed values (case 3 above), so we output 00 (header), then 111 (-1 on 3 bits signed) and finally 110 (-2 on 3 bits signed). At this point only zeros remain in the block, so 1010 (the ZEROS_TO_END header) is placed in the bitstream to conclude the coding of this block.

Coding of the second block from Figure 3 proceeds in the same way, however more data is required due to the presence of more nonzero value as well as the larger average magnitude of the values. As before, encoding proceeds until the last nonzero value in the scanned sequence (-1) is reached, after which a ZEROS_TO_END header is placed in the bitstream. When the last block of the downsampled V plane has been coded, zero padding ensures the file ends at the end of a 16-bit word (for compliance with the 16 bit wide memory we are using).

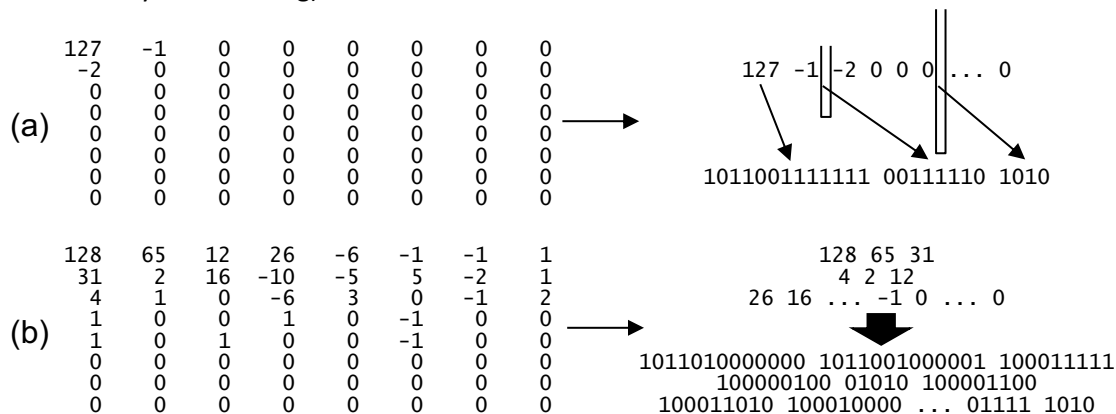


Figure 5. Lossless coding of each block from Figure 3.

3 Image Decompression

In this section, the whole decompression process for an image of size 16 x 16 is performed to illustrate the flow of data from one stage to another and elaborate on each step's internal details in the decompression. Since the aim of this project is the hardware implementation of the decompression specification, *during the decompression, we only use fixed point arithmetic*, which has a more straightforward hardware design and generally uses fewer logic resources compared to floating point.

3.1 Lossless Decoding and Dequantization (Milestone 3 in Hardware)

Our reference image “example.ppm” is shown in Figure 6. The encoded bitstream is given in Figure 7. The header is 8 bytes. The first 4 bytes must be “DEADBEEF” in hex, identifying a .mic18 file (future revisions may use a different identifier). After this, in the following 2 bytes, the most significant bit indicates which quantization matrix to use (in this example, bit 0 means use Q_0), and the 15 least significant bits indicate the width (bits 000 0000 0001 0000 indicate a width of 16). The following 2 bytes indicate the height, which is also 16 in our example. After the header, the main content of the image, the sequence of compressed blocks, begins.



Figure 6. The reference image example.ppm to be compressed.

```
DE EA BE EF 00 10 00 10 B3 A4 2A 0E 99 6C DA 0A
1C 60 99 DA 74 6A 12 73 A2 1E 9C 27 67 0E C6 DD
F7 6D F6 7B 9B BA 09 0F 9C 36 0D 05 7E A7 21 67
A6 F9 38 38 82 E9 3B 09 F1 83 F8 FE 2B EE D1 D5
9C E1 06 71 46 CE 0A 4D 9F 37 2A 0A C9 CE 82 20
89 D0 4E 53 F3 04 E3 2A F9 C1 C2 6D 07 06 54 E2
60 99 D3 67 5A 4C DC FC 24 2C 37 27 24 3D 39 09
38 08 76 36 EF BB 6E B4 14 5A 33 0E 99 C2 E0 A6
07 21 90 58 23 67 53 A3 3D BF 61 07 01 D9 7E 31
97 D5 A3 E3 F3 68 74 1E 23 2D 83 CE 8C 93 39 29
1A 82 07 8F 9C 1C 1C E4 3F 47 09 D0 42 6A CF D0
A4 CD B3 34 9B 64 B8 C2 08 72 13 04 17 DA 74 1E
70 7E 9A B3 C4 DA 0B 3A 74 E6 0E 55 93 85 E8 FC
C0 3F C7 07 34 00
```

Figure 7. In the compressed bitstream example.mic18, the order of the bytes is across first and then down.

Figure 8 depicts the block decoding process. Variable length prefix coding is used for the header, which means depending on the first few bits of the header, we may need to read additional bits to determine the remainder of the header. We start by reading 2 bits: if we get 00, then two 3-bit values follow; if we get 01, then one 3-bit value follows; if we get 11, then a run of zeros follows; and if we get 10, then we need to read more bits to complete the header. In the case where the first two bits are 10, if the next bit we read is 0, then one 6-bit value follows; otherwise, we need to read one more bit to complete the header. If this bit is 1, then one 9-bit value follows; otherwise, there is a run of zeros to the end of the block. Following this procedure, values are decoded one after another and placed in the 8x8 block according to the scan pattern of Figure 4. A block is complete either when a ZEROS_TO_END header is decoded (in which case the rest of the block is filled with 0s), or when 64 values (including runs of zeros) have been decoded, making the block complete. Note that the values are signed (2's complement format).

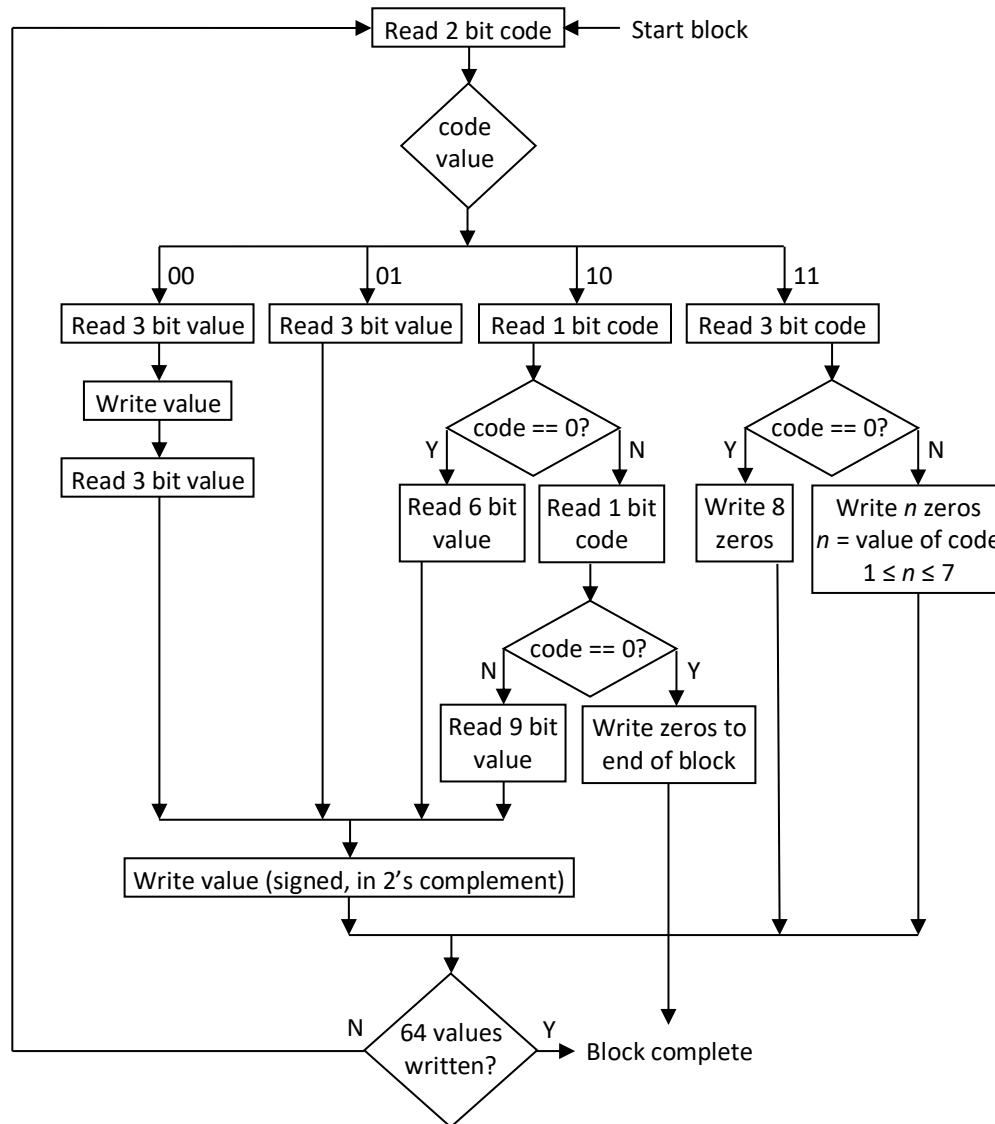


Figure 8. Lossless decoding of a block based on the prefix code.

The sequence begins with hex “B3 A4 2A ...” which is “101100111010010000101010 ...” in binary, and when applied to this bitstream, the result of the lossless decoding procedure is as shown on the following two pages in Figures 9 and 10.

The Y blocks appear first in the bitstream, row by row. Y 0,0 contains Y data for pixel rows 0-7 and columns 0-7 of the Y plane from the downsampled image (equivalent to the Y plane from the converted image since Y is not downsampled). Similarly, Y 0,1 contains rows 0-7, columns 8-15; Y 1,0 contains rows 8-15, columns 0-7; Y 1,1 contains rows 8-15, columns 8-15.

Y 0 0 -----									
1011001110100 : 116	11010 : 0 0								
100001010 : 10	00100001 : -4 1								
100000111 : 7	11101 : 0 0 0 0 0								
01001 : 1	00111000 : -1 0								
100101101 : -19	01001 : 1	116	10	-10	5	0	0	0	0
100110110 : -10	11011 : 0 0 0	7	-19	1	1	-2	0	0	0
100000101 : 5	00111000 : -1 0	1	-2	-5	-3	1	0	0	0
00001110 : 1 -2	01110 : -2	-2	1	1	-4	-1	0	0	0
00110000 : -2 0	11000 : 0 0 0 0 0 0 0 0	0	1	0	0	-2	0	0	0
01001 : 1	11011 : 0 0 0	-1	0	1	0	0	-2	0	0
100111011 : -5	01110 : -2	1	0	-1	0	0	0	-2	0
01001 : 1	11111 : 0 0 0 0 0 0 0	0	0	0	0	0	0	0	-1
11010 : 0 0	01110 : -2								
00110101 : -2 -3	11011 : 0 0 0								
00001001 : 1 1	01111 : -1								
00111001 : -1 1									
(a)									
Y 0 1 -----									
1011001111011 : 123	00111000 : -1 0								
100110111 : -9	100000101 : 5								
01110 : -2	11010 : 0 0								
100000100 : 4	01001 : 1								
100001111 : 15	11011 : 0 0 0	123	-9	-8	-3	3	-1	1	0
100111000 : -8	00001001 : 1 1	-2	15	6	1	-3	0	0	0
01101 : -3	11110 : 0 0 0 0 0 0	4	5	-7	-2	0	0	-1	0
100000110 : 6	00110000 : -2 0	-1	0	-1	5	1	0	0	0
100000101 : 5	01111 : -1	0	0	0	1	-2	0	0	0
01111 : -1	11100 : 0 0 0 0	-1	-1	0	0	-1	2	0	0
11010 : 0 0	01111 : -1	0	0	0	0	0	0	-2	0
100111001 : -7	11100 : 0 0 0 0	0	0	0	0	0	0	-1	2
00001011 : 1 3	01010 : 2								
00111101 : -1 -3	11111 : 0 0 0 0 0 0 0								
00110111 : -2 -1	01110 : -2								
11001 : 0	11010 : 0 0								
00111000 : -1 0	00111010 : -1 2								
(b)									
Y 1 0 -----									
1011001110011 : 115	100000100 : 4								
100001000 : 8	01001 : 1								
00110011 : -2 3	11010 : 0 0								
100010100 : 20	00001001 : 1 1								
01101 : -3	11001 : 0	115	8	-3	-8	2	1	0	1
100111000 : -8	01001 : 1	-2	20	-3	-3	2	0	1	0
00101001 : -3 1	11111 : 0 0 0 0 0 0 0	3	1	-9	0	1	0	1	0
00110110 : -2 -2	00110000 : -2 0	-2	-1	-1	4	1	0	0	0
01111 : -1	01001 : 1	-2	1	0	0	-2	0	0	0
100110111 : -9	11000 : 0 0 0 0 0 0 0 0	0	1	0	0	0	2	0	0
00101010 : -3 2	11001 : 0	0	0	0	0	0	0	-1	-1
00001010 : 1 2	01010 : 2	0	0	0	0	0	0	0	1
11001 : 0	11111 : 0 0 0 0 0 0 0								
00111001 : -1 1	00111000 : -1 0								
11010 : 0 0	00111000 : -1 0								
00001000 : 1 0	01001 : 1								
(c)									
Y 1 1 -----									
1011010000011 : 131	11001 : 0								
100000110 : 6	00111001 : -1 1								
01010 : 2	00100001 : -4 1								
100111000 : -8	11101 : 0 0 0 0 0								
100110000 : -16	00111001 : -1 1	131	6	1	-6	-1	1	0	0
01001 : 1	00001001 : 1 1	2	-16	-3	-1	1	0	0	0
100111010 : -6	00111000 : -1 0	-8	-6	-9	1	1	0	0	0
01101 : -3	00001000 : 1 0	0	1	3	-4	-1	0	0	0
100111010 : -6	01110 : -2	0	1	1	1	-2	0	0	0
11010 : 0 0	11000 : 0 0 0 0 0 0 0 0	-3	-1	1	0	0	-2	0	0
01001 : 1	11011 : 0 0 0	0	1	1	0	0	0	-2	0
100110111 : -9	01110 : -2	-1	0	0	0	0	0	0	-2
00111111 : -1 -1	11111 : 0 0 0 0 0 0 0								
00001001 : 1 1	01110 : -2								
00001011 : 1 3	11011 : 0 0 0								
00001101 : 1 -3	01110 : -2								
(d)									

Figure 9. Lossless decoding of the Y blocks in example.mic18.

Next in the bitstream comes the data for the U and V planes of the downsampled image, recall that they have half the columns of the Y plane due to downsampling. In the same way as for Y, U 0,0 contains U data for pixel rows 0-7, columns 0-7 of the downsampled image; U 1,0 contains rows 8-15, columns 0-7; V 0,0 contains V data for rows 0-7, columns 0-7; V 1,0 contains rows 8-15, columns 0-7.

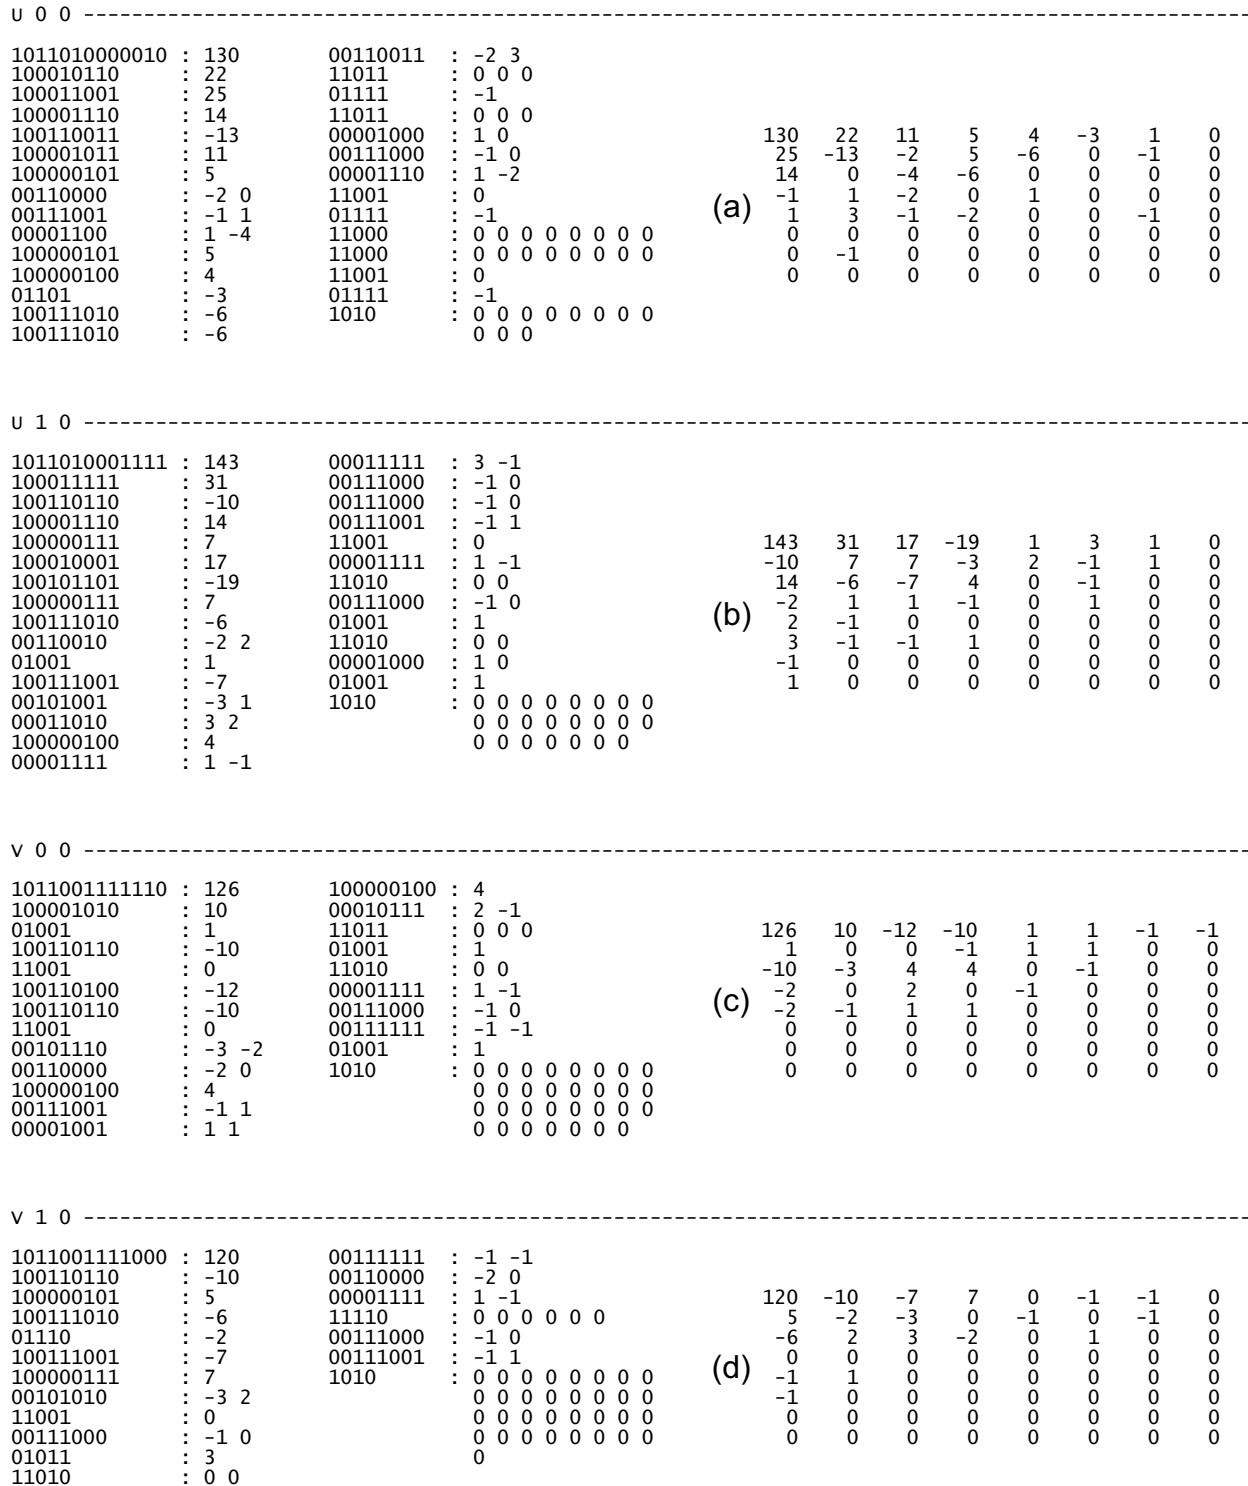


Figure 10. Lossless decoding of the downsampled U and V blocks in example.mic18.

Upon completion of decoding, we have 8 blocks of data, 4 for Y (16/8 x 16/8), 2 for U and 2 for V. The matrices in the rightmost column of Figures 9 and 10 collectively make up the quantized image. From the flow illustrated in Figure 1, the next step is dequantization, which leads to the transformed image. This is done using pointwise multiplication between the losslessly decoded data (matrices in the rightmost column above) and the appropriate quantization matrix (Q_0 in this case). For matrices A , B and C , pointwise (or element-wise) multiplication is defined in Equation 6 below.

$$A = B .* C \iff [A]_{i,j} = [B]_{i,j} * [C]_{i,j} \quad (6)$$

The quantized images are shown on the leftmost column on page 12 and the result of dequantization is shown in the central column.

3.2 IDCT (Inverse Discrete Cosine Transform) (Milestone 2 in Hardware)

After dequantizing the blocks, the IDCT must be performed to recover the downsampled image. The IDCT is defined in Equation 7:

$$S = (int) \frac{1}{65536} \left(C^T \times (int) \frac{1}{256} (S' C) \right) \quad (7)$$

where the matrix S' is the dequantized 8x8 block, and the output of the transform S is the block of 8x8 samples from the downsampled image. The element on the i^{th} row and j^{th} column of the 8x8 matrix C is:

$$[C]_{i,j} = \begin{cases} (int) \left(\sqrt{\frac{1}{8}} \cos\left(\frac{i\pi}{8}(j + 0.5)\right) \times 4096 \right) & \text{if } i = 0 \\ (int) \left(\sqrt{\frac{2}{8}} \cos\left(\frac{i\pi}{8}(j + 0.5)\right) \times 4096 \right) & \text{otherwise} \end{cases} \quad (8)$$

$$C = \begin{bmatrix} 1448 & 1448 & 1448 & 1448 & 1448 & 1448 & 1448 & 1448 \\ 2008 & 1702 & 1137 & 399 & -399 & -1137 & -1702 & -2008 \\ 1892 & 783 & -783 & -1892 & -1892 & -783 & 783 & 1892 \\ 1702 & -399 & -2008 & -1137 & 1137 & 2008 & 399 & -1702 \\ 1448 & -1448 & -1448 & 1448 & 1448 & -1448 & -1448 & 1448 \\ 1137 & -2008 & 399 & 1702 & -1702 & -399 & 2008 & -1137 \\ 783 & -1892 & 1892 & -783 & -783 & 1892 & -1892 & 783 \\ 399 & -1137 & 1702 & -2008 & 2008 & -1702 & 1137 & -399 \end{bmatrix} \quad (9)$$

Evaluating all cases in Equation 8 produces Equation 9. Equation 7 is defined similarly to the DCT from Equation 2. However, the important differences are the intermediate and final scalar divisions and the use of fixed-point coefficients. To reiterate, the reason for choosing fixed-point (i.e. integer) operations is that they facilitate a more efficient hardware implementation than floating-point operations.

Fixed-point, however, brings in necessary approximations that we achieve by using an intermediate bit-width larger than the one necessary to store the IDCT outputs. For example, the cosine coefficients (which are between -1 and 1) are left-shifted by 12 positions (i.e. multiplied by 4096 as shown in Equation 7) to obtain an integer-only fixed-point value before processing (-4096 to 4096). Because the sample matrix is multiplied on both sides by the coefficient matrix, the data will be left-shifted by 24 positions in total. The first adjustment of a right shift by 8 positions is provided after the post-multiplication $S' C$ and the second adjustment is given by a right shift by 16 positions after the pre-multiplication $C^T (S' C)$. The reason why the first adjustment is only on 8 bits is because it retains more precision, and thus, it is more immune to the truncation errors.

Finally, a clipping function is used to ensure that rounding and overflows that have led to values out of the 8-bit unsigned range (0 to 255) are addressed by saturation (i.e. values smaller than 0 are forced to 0 and values greater than 255 are forced to 255). The results of dequantization and IDCT are as follows:

Y 0,0																								
116	10	-10	5	0	0	0	0	928	40	-80	40	0	0	0	0	7	142	145	128	133	136	140	127	
7	-19	1	1	-2	0	0	0	28	-152	8	16	-32	0	0	0	131	8	122	139	129	151	134	131	
1	-2	-5	-3	1	0	0	0	8	-16	-80	-48	32	0	0	0	122	116	23	152	195	157	139	97	
-2	1	1	-4	-1	0	0	0	-16	16	16	-128	-32	0	0	0	134	139	111	17	131	163	112	89	
0	1	0	0	-2	0	0	0	0	16	0	0	-128	0	0	0	133	128	150	121	36	139	124	64	
-1	0	1	0	0	-2	0	0	-16	0	32	0	0	-128	0	0	115	140	135	167	146	33	107	79	
1	0	-1	0	0	0	-2	0	32	0	-64	0	0	0	-128	0	146	120	122	107	120	96	0	127	
0	0	0	0	0	0	0	-1	0	0	0	0	0	0	0	-64	133	122	136	176	158	116	101	3	
Y 0,1																								
123	-9	-8	-3	3	-1	1	0	984	-36	-64	-24	48	-16	32	0	130	154	125	134	145	139	156	5	
-2	15	6	1	-3	0	0	0	-8	120	48	16	-48	0	0	0	148	134	153	152	143	129	0	163	
4	5	-7	-2	0	0	-1	0	32	40	-112	-32	0	0	-64	0	103	138	147	135	130	10	149	152	
-1	0	-1	5	1	0	0	0	-8	0	-16	160	32	0	0	0	108	85	159	109	1	195	119	147	
0	0	0	1	-2	0	0	0	0	0	0	32	-128	0	0	0	139	96	109	33	182	144	110	139	
-1	-1	0	0	-1	2	0	0	-16	-32	0	0	-64	128	0	0	133	82	18	164	170	153	89	153	
0	0	0	0	0	0	-2	0	0	0	0	0	0	0	-128	0	97	7	144	168	160	126	138	142	
0	0	0	0	0	0	-1	2	0	0	0	0	0	0	0	128	12	145	153	189	172	112	118	153	
Y 1,0																								
115	8	-3	-8	2	1	0	1	920	32	-24	-64	32	16	0	32	142	94	171	173	128	109	54	13	
-2	20	-3	-3	2	0	1	0	-8	160	-24	-48	32	0	32	0	128	111	175	141	169	80	8	143	
3	1	-9	0	1	0	1	0	24	8	-144	0	32	0	64	0	131	146	154	168	95	14	99	143	
-2	-1	-1	4	1	0	0	0	-16	-16	-16	128	32	0	0	0	126	144	155	114	13	63	146	101	
-2	1	0	0	-2	0	0	0	-32	16	0	0	-128	0	0	0	102	139	104	15	122	96	143	122	
0	1	0	0	0	2	0	0	0	32	0	0	0	128	0	0	124	111	8	136	137	128	132	143	
0	0	0	0	0	0	-1	-1	0	0	0	0	0	0	-64	-64	139	14	134	139	154	107	153	125	
0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	64	15	120	150	127	127	144	127	135	
Y 1,1																								
131	6	1	-6	-1	1	0	0	1048	24	8	-48	-16	16	0	0	19	133	126	120	113	121	141	152	
2	-16	-3	-1	1	0	0	0	8	-128	-24	-16	16	0	0	0	110	7	179	202	181	151	140	141	
-8	-6	-9	1	1	0	0	0	-64	-48	-144	16	32	0	0	0	141	149	37	163	136	120	134	152	
0	1	3	-4	-1	0	0	0	0	16	48	-128	-32	0	0	0	157	164	169	29	141	147	157	154	
0	1	1	1	-2	0	0	0	0	16	32	32	-128	0	0	0	173	190	186	152	4	147	153	132	
-3	-1	1	0	0	-2	0	0	-48	-32	32	0	0	-128	0	0	178	193	182	135	108	6	162	160	
0	1	1	0	0	0	-2	0	0	32	64	0	0	0	0	0	98	115	135	135	127	135	18	145	
-1	0	0	0	0	0	0	-2	-32	0	0	0	0	0	-128	-128	152	132	131	144	145	139	139	13	
U 0,0																								
130	22	11	5	4	-3	1	0	1040	88	88	40	64	-48	32	0	131	185	168	178	185	159	168	154	
25	-13	-2	5	-6	0	-1	0	100	-104	-16	80	-96	0	-32	0	168	149	130	104	149	193	156	150	
14	0	-4	-6	0	0	0	0	112	0	-64	-96	0	0	0	0	168	124	87	57	121	187	154	150	
-1	1	-2	0	1	0	0	0	-8	16	-32	0	32	0	0	0	175	154	86	80	111	110	113	110	
1	3	-1	-2	0	0	-1	0	16	48	-32	-64	0	0	-64	0	182	140	86	97	114	79	72	101	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	187	94	89	91	116	105	54	133	
0	-1	0	0	0	0	0	0	0	-32	0	0	0	0	0	0	193	122	138	123	114	86	41	143	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	159	167	181	174	121	46	61	160	
U 1,0																								
143	31	17	-19	1	3	1	0	1144	124	136	-152	16	48	32	0	188	158	194	173	130	126	128	176	
-10	7	7	-3	2	-1	1	0	-40	56	56	-48	32	-32	32	0	174	174	162	136	78	62	83	177	
14	-6	-7	4	0	-1	0	0	112	-48	-112	64	0	-32	0	0	177	171	165	113	83	77	90	172	
-2	1	1	-1	0	1	0	0	-16	16	16	-32	0	64	0	0	183	159	159	116	68	79	119	174	
2	-1	0	0	0	0	0	0	32	-16	0	0	0	0	0	0	186	161	191	137	16	29	130	180	
3	-1	-1	1	0	0	0	0	48	-32	-32	64	0	0	0	0	163	132	168	141	37	53	134	137	
-1	0	0	0	0	0	0	0	-32	0	0	0	0	0	0	0	173	155	172	182	160	171	184	147	
1	0	0	0	0	0	0	0	32	0	0	0	0	0	0	0	139	167	172	174	180	174	165	157	
V 0,0																								
126	10	-12	-10	1	1	-1	-1	1008	40	-96	-80	16	16	-32	-32	115	114	90	106	93	105	117	113	
1	0	0	-1	1	1	0	0	4	0	0	-16	16	32	0	0	108	125	139	171	128	104	112	116	
-10	-3	4	4	0	-1	0	0	-80	-24	64	64	0	-32	0	0	104	129	169	215	151	106	112	118	
-2	0	2	0	-1	0	0	0	-16	0	32	0	-32	0	0	0	107	129	164	203	142	110	120	117	
-2	-1	1	1	0	0	0	0	-32	-16	32	32	0	0	0	0	104	139	162	180	126	113	124	114	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	97	151	165	170	126	114	118	113	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	99	142	133	145	130	117	106	112	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	108	122	84	108	130	120	98	110	
V 1,0																								
120	-10	-7	7	0	-1	-1	0	960	-40	-56	56	0	-16	-32	0	99	116	100	120	117	108	134	99	
5	-2	-3	0	-1	0	-1	0	20	-16	-24	0	-16	0	-32	0	102	122	106	130	142	132	145	113	
-6	2	3	-2	0	1	0	0	-48	16	48	-32	0	32	0	0	98	122	104	130	160	147	138	110	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	100	122	102	127	166	150	123	100	
-1	1	0	0	0	0	0	0	-16	16	0	0	0	0	0	0	109	123	107	128	167	153	122	106	
-1	0	0	0	0	0	0	0	-16	0	0	0	0	0	0	0	113	114	106	122	150	144	123	116	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	116	103	102	110	119	121	115	114	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	126	103	105	106	100	104	108	111	

3.3 Upsampling and Colourspace Conversion (Milestone 1 in Hardware)

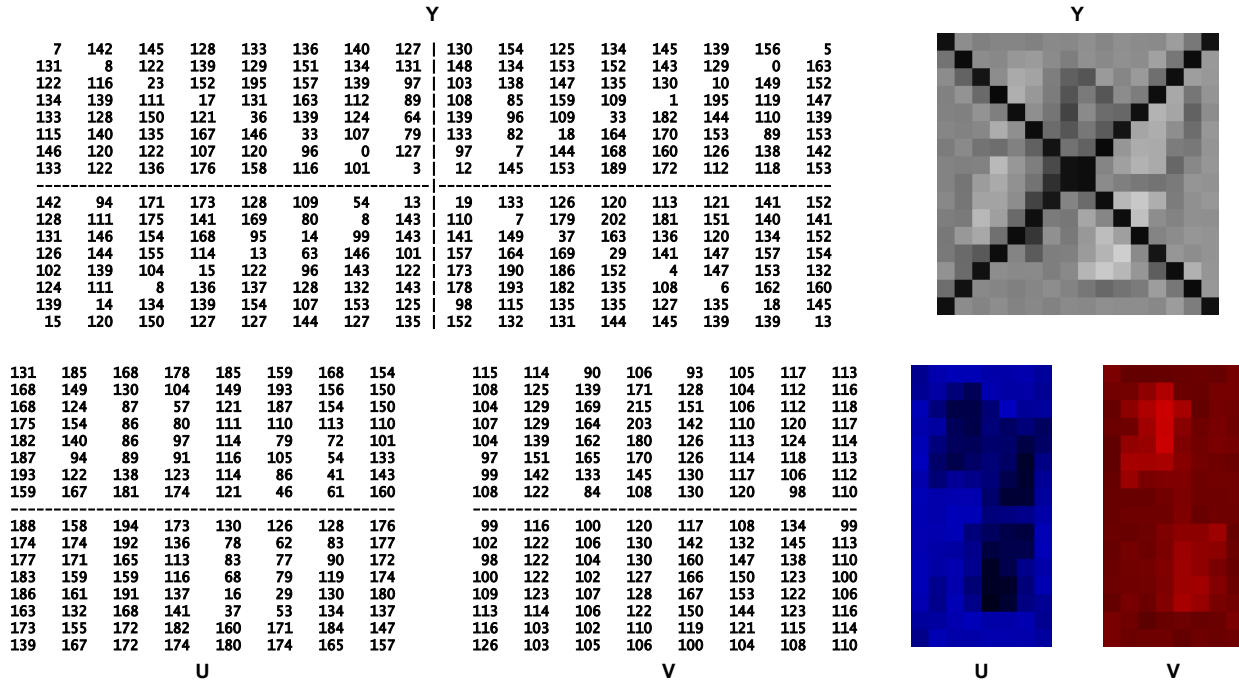


Figure 12. The output of IDCT is the downsampled YUV data.

The output of IDCT is downsampled YUV data, where U and V are horizontally downsampled by 2, as shown in Figure 12. The downsampled data becomes the even columns in the upsampled image, and the odd columns are generated by horizontal interpolation. For each row (rows are processed independently), we interpolate with a filter known as a finite impulse response (FIR) filter, which uses a weighted sum of a finite number of samples to produce each reconstructed sample. In our case, that finite number is 6, so the filter is called a 6-tap filter. Given n samples in each row of the U and V planes, we produce $2n$ samples for the corresponding row of the reconstructed plane as follows:

$$U'[j] = \begin{cases} U[\frac{j}{2}] & j \text{ even} \\ (int)\frac{1}{256}(21U[\frac{j-5}{2}] - 52U[\frac{j-3}{2}] + 159U[\frac{j-1}{2}] + 159U[\frac{j+1}{2}] - 52U[\frac{j+3}{2}] + 21U[\frac{j+5}{2}] + 128) & j \text{ odd} \end{cases} \quad (10)$$

where U is the downsampled version and U' is the upsampled version. An analogous definition applies to V. In our example, with an image of size 16x16, n is 8. In the project, the image will be of size 320x240, so n will be $320/2 = 160$. The borders of **each row** are handled by replicating the respective border sample, i.e. $U[0] = U[-1] = U[-2]$, and $U[n-1] = U[n] = U[n+1] = U[n+2]$. Note again the use of integer operations. These coefficients are borrowed from the ISO/IEC 13818-2 document on MPEG2 digital video decoding. Applying this filter to the U and V planes shows the result in Figure 13.

Having restored the complete Y, U and V planes, the only remaining step is the colourspace conversion from YUV back to RGB. As discussed above (Equation 1), the conversion from RGB to YUV is a matrix multiplication. In essence, the RGB values of a given pixel are viewed as a vector in 3-space, and the matrix multiplication is a change of basis yielding another vector in 3-space. Since this is a linear transformation and the determinant of the matrix is nonzero, we may find its inverse, which will




Figure 13. The result after horizontally upsampling U and V.

provide the transformation back to the RGB basis from the YUV basis. This is applied to each pixel independently. Accounting for the vector addition to first shift the basis, the inverse conversion of the one above is:

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = (int) \frac{1}{65536} \left(\begin{bmatrix} 76284 & 0 & 104595 \\ 76284 & -25624 & -53281 \\ 76284 & 132251 & 0 \end{bmatrix} \left(\begin{bmatrix} Y \\ U \\ V \end{bmatrix} - \begin{bmatrix} 16 \\ 128 \\ 128 \end{bmatrix} \right) \right) \quad (11)$$

As in the case of IDCT, an identical clipping function is used to deal with the under/overflow errors. The outputs R, G, and B are each on 8 bits unsigned. The final result is shown in Figure 14.




Figure 14. After colourspace conversion, we obtain the final decompressed RGB image.

The final decompressed image, as shown on the left (a) in Figure 15, requires $16 \times 16 \times 3 = 768$ bytes (plus a header of 13 bytes), whereas the compressed image required only 214 bytes, giving a compression ratio of more than 3x. The original image appears on the right (c). As you can see, there is some reduction in the quality due to information lost during the compression and decompression process. The center image (b) shows the decompressed image for the case when the quantization matrix Q_1 from Equation 3 is used (instead of Q_0 as is the case for this example, image (a)), which during compression retains more frequency information for use during reconstruction. The use of this matrix improves the quality but reduces the compression. This discussion is summarized in the signal-to-noise ratio (SNR) values and file sizes given for both cases in Table 1.

	Q_0 image	Q_1 image	Original image
SNR (dB)	18.65	20.40	∞
File size (bytes, including header)	214	356	781
Compression ratio	3.65	2.19	1

Table 1. Quality versus compression tradeoff for the images using quantization matrices Q_0 and Q_1 .




Figure 15. Final decompressed images using quantization matrices Q_0 (a) and Q_1 (b) and the original image (c).

Equation 12 defines the signal-to-noise ratio where N is the number of pixels, $error$ is the root mean squared error, and the 255 on the numerator is the maximum value of any pixel:

$$SNR (dB) = 20 \log_{10} \left(\frac{255}{\sqrt{\frac{error}{N}}} \right) \text{ where } error = \sum_{n=0}^{N-1} (reference_pixel_n - image_pixel_n)^2 \quad (12)$$

This part of the document has provided the barebones of image compression (sufficient to work on the project). Further information on image compression can be found in the following references:

- [1] JPEG: Still Image Data Compression Standard, by William B., Pennebaker, Joan L., Mitchell, Springer; 1 edition (2006), ISBN: 0442012721
- [2] Digital Video and HDTV Algorithms and Interfaces, by Charles Poynton, Publisher: Morgan Kaufmann (2003), ISBN: 1558607927
- [3] Video Codec Design: Developing Image and Video Compression Systems, by Iain Richardson, Publisher: John Wiley & Sons (2002), ISBN: 0471485535
- [4] Introduction to Data Compression, Third Edition, by Khalid Sayood, Publisher: Morgan Kaufmann (2005), ISBN: 012620862X
- [5] Image and Video Compression for Multimedia Engineering: Fundamentals, Algorithms, and Standards, by Yun Q. Shi, Huifang Sun, Publisher: CRC Press (1999), ISBN: 0849334918
- [6] Digital Video Compression, by Peter Symes, Publisher: McGraw-Hill (2003), ISBN: 0071424873

4 Project Guidelines

The design task of implementing the overall decoder has been partitioned into three milestones, each of which can be completed within one week. As part of the validation methodology, implementation will proceed from the VGA display backward to the compressed input. Starting from the display end enables a visual progress check at each milestone, providing the framework for incremental integration of each newly designed module into the overall system.

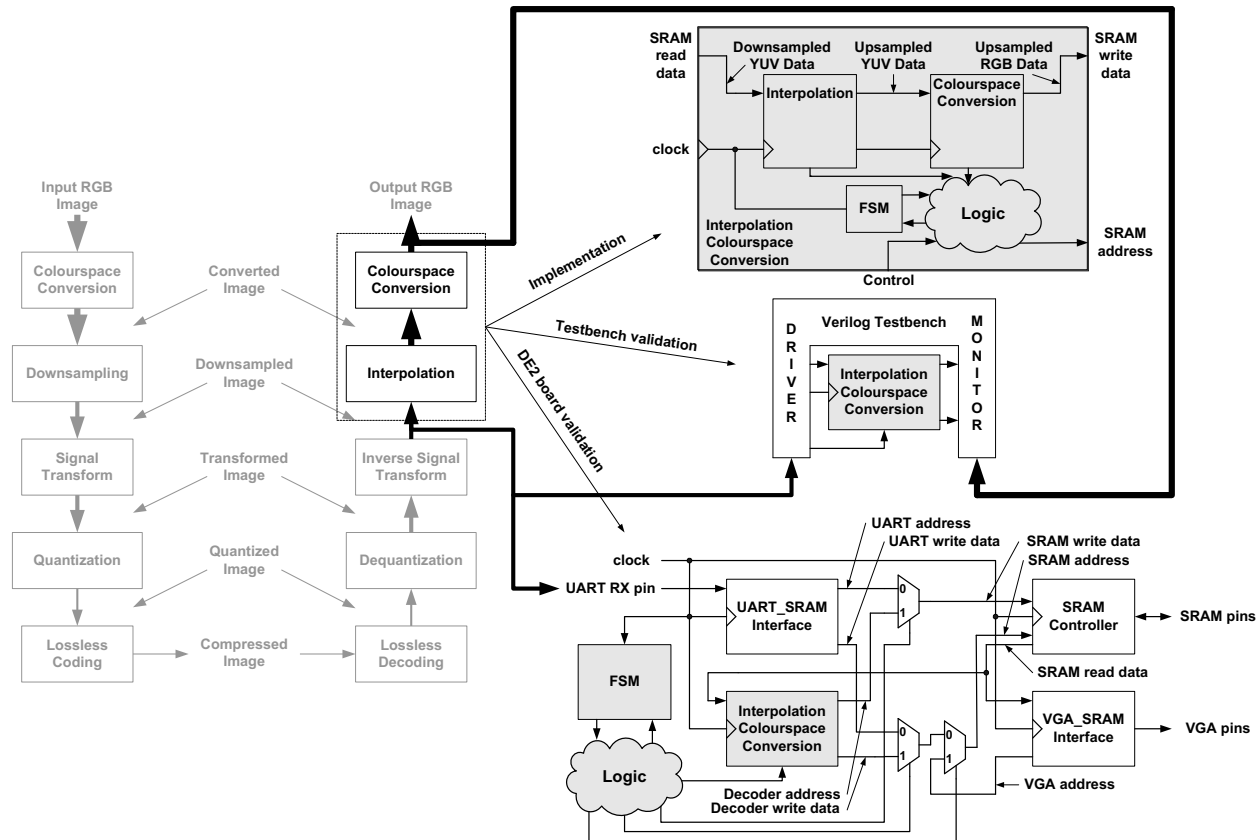


Figure 16. Overview of the upsampling and colourspace conversion hardware implementation (milestone 1).

4.1 The Software Model and its Integration in the Design Process

Starting with a .mic18 file, there are 3 separate source codes, 1 for each milestone. The input/output files are organized exactly like the SRAM in hardware, that is $2^{18} = 262144$ locations with 16 bits per location. Each source code reads data from a software model of the SRAM, performs the necessary computation, and then writes the data back to the SRAM software model. The decoding process starts at milestone 3, and the source code "decode_m3" takes a .mic18 file and produces a .sram_d2 file. Continuing into milestone 2, the source code "decode_m2" takes a .sram_d2 file and produces a .sram_d1 file. Finally, for milestone 1, "decode_m1" takes a .sram_d1 file and produces a .sram_d0 file and a .ppm file.

By separating the decoding process, we can do part of the decoding in software and part of it in hardware. The decoding software handles the first part of the decoding task and the intermediate data can be passed to a testbench or to the board, where the hardware can take over the remainder of the

decoding task to produce the image. In the case of the testbench, detailed feedback can be provided, which aids in tracing implementation errors, while the board provides less feedback but a full practical operational test. For example, to validate if milestone 1 is working properly in hardware, the .sram_d1 file can be used to initialize the SRAM, and after the hardware simulation is finished, the .sram_d0 file can be used to validate if the contents of SRAM are correct.

4.2 Milestone 1: Upsampling and Colourspace Conversion

Figure 16 shows the implementation and validation plan for milestone 1. In this first week, colourspace conversion and interpolation will be tackled as indicated in the portion of the figure which shows the flow from Figure 1. This will also require some modifications to the existing state machines, to grant the decoder access to the external memory (by multiplexing). In fact, all the shaded portions of the figure will have to be built or modified. The SRAM Controller, UART SRAM Interface and VGA SRAM Interface shown in the figure have been provided in lab 5.

Figure 17 shows the memory layout which is used for milestone 1. The last portion of the memory holds the completed RGB image stored in {R0 G0}, {B0 R1}, {G1 B1} ... format (as in lab 5 experiment 1). This is where the decompressed image will reside for all three milestones and, thus, for the completed project as well. Where it differs from the completed project is in the fact that downsampled YUV data is stored in the SRAM as decoder stimuli, since the earlier parts of the decoder are not yet implemented. By default (from the backbone in lab 5), the UART SRAM interface writes data in the SRAM starting at address 0 and the VGA SRAM interface starts reading from address 0. Clearly, the VGA SRAM interface must be modified to read the RGB data from the correct segment, as indicated in Figure 17. *Also, the UART receiver in lab 5 experiment 4 strips the header from the .ppm file, which is no longer needed.*

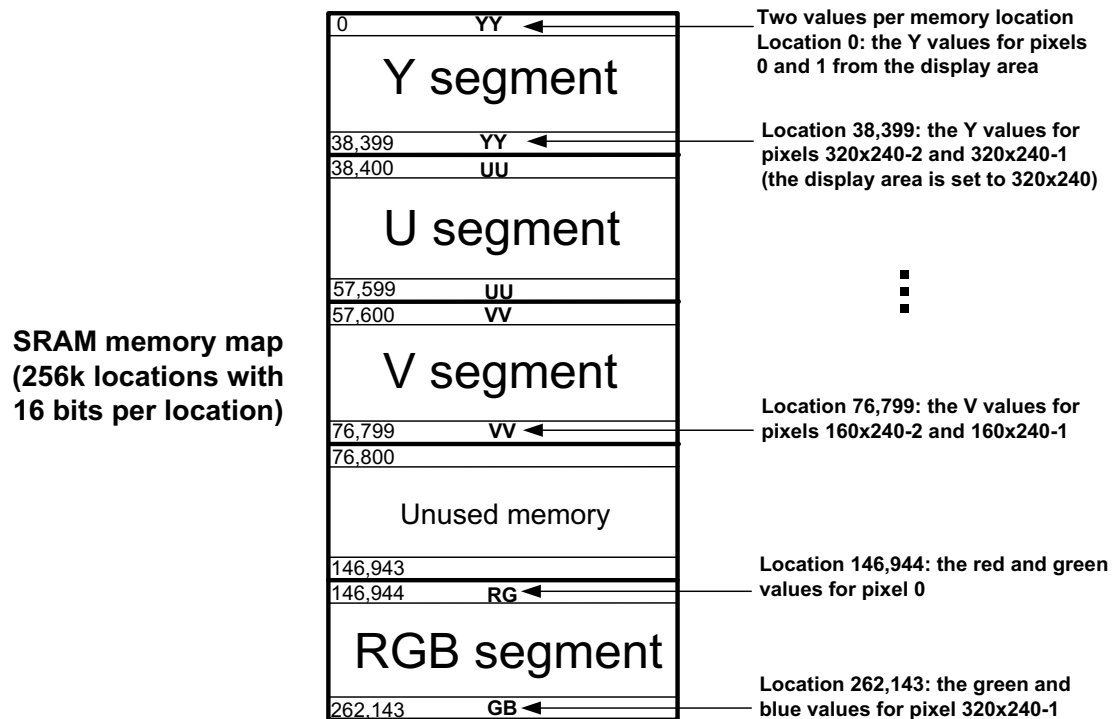


Figure 17. Memory layout for milestone 1.




Figure 18. Instantiation of a single 32-bit multiplier.

Figure 18 shows how a single 32-bit integer multiplier can be instantiated in your design. As will be elaborated on soon, to end up with a resource-efficient implementation, multipliers must be shared because they are costly in terms of logic resources. Specifically, for this milestone, you are given the **constraint of 3 multipliers (with input operands on 32 bits or less)**. **Note that these 3 multipliers are shared among colourspace conversion and upsampling, including the symmetry in upsampling as summarized in Equation (13).** Thus, there will be three instantiations of the form for Figure 18 (note that you are required to meet utilization constraints). Each of these tasks comprises more than one multiplication, so the multipliers for each task must be shared. For this reason, the instantiation given in Figure 18 should be used, and the signals Mult_op_1, Mult_op_2 and Mult_result should be multiplexed.

By explicitly instantiating multipliers, we are certain to have that many multipliers used in hardware. For example, in the following code snippet:

```

always_comb begin
    if (select == 1'b0) a = b*c;
    else a = d*e;
end

```

The multiplications happen at mutually exclusive times, so we could have 1 hardware multiplier with b and d multiplexed at one input to the multiplier, and c and e multiplexed at the other input to the multiplier. However, this not explicitly stated, so the computer-aided design (CAD) tool (in our case, Quartus) may choose to instantiate 2 hardware multipliers to compute both $b*c$ and $d*e$, and then take these 2 results and multiplex them to obtain the output a . To avoid this problem, we can explicitly specify that 1 (and only 1) multiplier should be used by multiplexing the operands:

```

always_comb begin
    if (select == 1'b0) begin
        op1 = b;
        op2 = c;
    end
    else begin
        op1 = d;
        op2 = e;
    end
end
assign a = op1*op2;

```

Next, some architectural examples are provided regarding how the multipliers can be shared.

Figure 19 shows a straightforward implementation of a FIR filter, where the samples are stored in a shift-register structure. In this case, each sample undergoes a constant multiplication, and the results are all combined using an adder tree. Although this is attractive from the throughput perspective, as one filtered sample may be obtained every clock cycle, it is costly from the resource standpoint, especially considering that one such filter would be required for each component, which must be upsampled (U and V). Furthermore, we must consider if we can actually receive a new U and a new V sample on every clock cycle; if we cannot, then on some clock cycles, this hardware will be under-utilized. A similar argument applies to whether or not we can process one upsampled U and one upsampled V per clock cycle after the FIR filter.

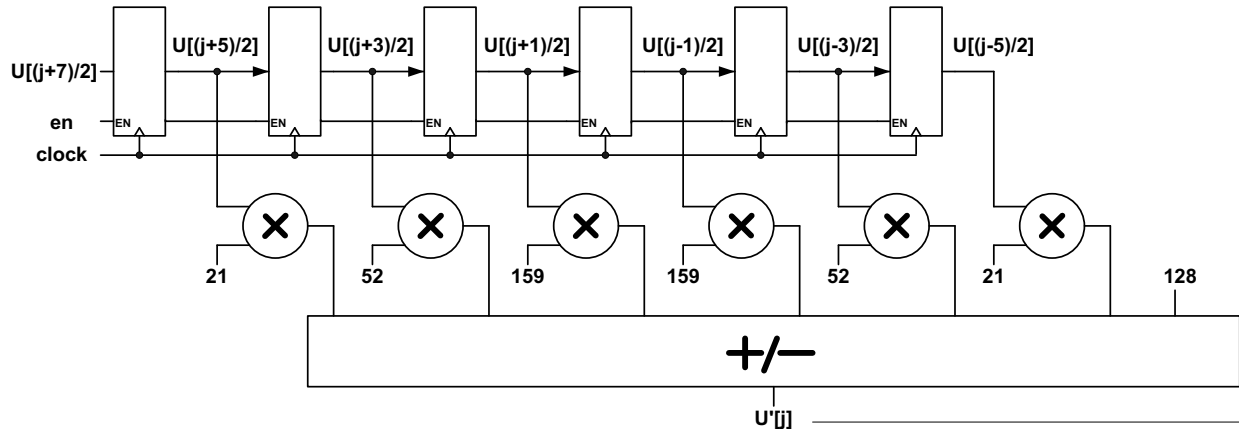


Figure 19. Resource inefficient FIR filter implementation.

Figure 20 shows how a single multiplier can be shared to accomplish the same computation as Figure 19, but at the expense of throughput. In this case, the boxed portion is known as a Multiply-Accumulate (MAC) unit, which works by always multiplying the two inputs, adding the result to the register (known as the accumulation register or accumulator), and then storing that result back in the accumulator, to become an operand in the next clock cycle. The FSM guides the selection of both the sample and the constant factor, as well as initialization, which in this case is to account for the +128 in the upsampling equations of Equation 10, also shown in Figure 19. Using this architecture, only one multiplier is required; however, each reconstructed sample requires 7 clock cycles to produce instead of 1, as in the previous case. This is an example of a typical architectural tradeoff between resource usage and throughput.

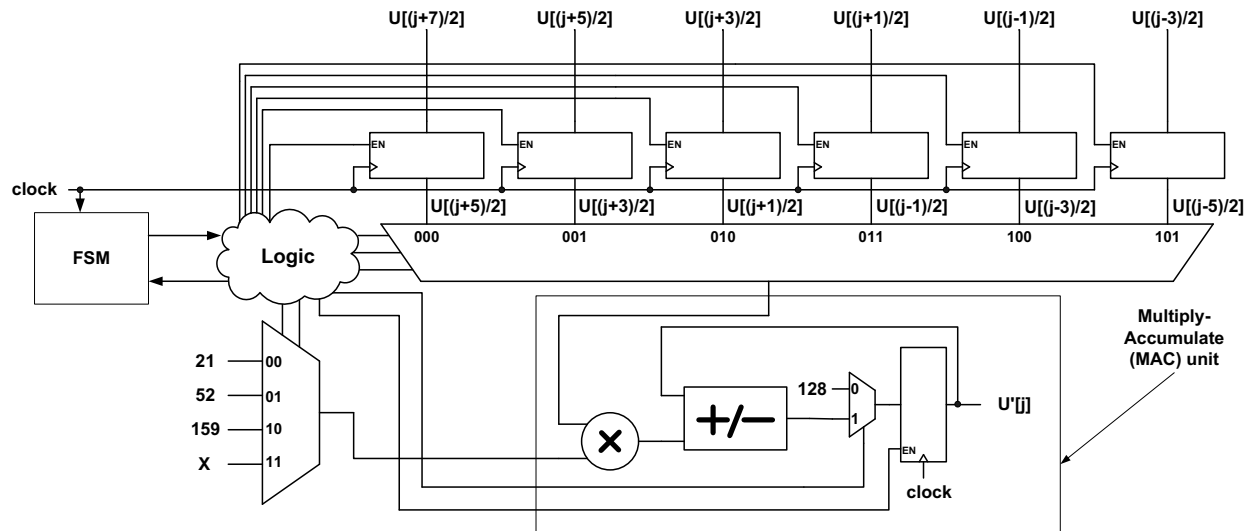


Figure 20. Resource-efficient FIR filter implementation.

Figure 21 shows a cost-effective FIR implementation similar to Figure 20, but with the important distinction that the multiplexer that selects between the different samples (as set up by the FSM) is eliminated. The same result can be obtained by leveraging the shift structure that is already present due to the reuse of the majority of the samples in calculating adjacent reconstructed samples. Now, instead of the FSM setting up select lines on a mux to obtain a desired sample, the calculation proceeds by holding the shift for the first clock cycle (during initialization of the accumulator to 128), and during the next 6 clock cycles, the samples are circularly shifted, each time the contents of the rightmost register being accumulated. At the end of these 6 clock cycles, each data returns to its original register, and the accumulation is complete (a new reconstructed sample has been produced). Although identical to Figure 20 in throughput, eliminating the large sample selector mux makes this design substantially more efficient.

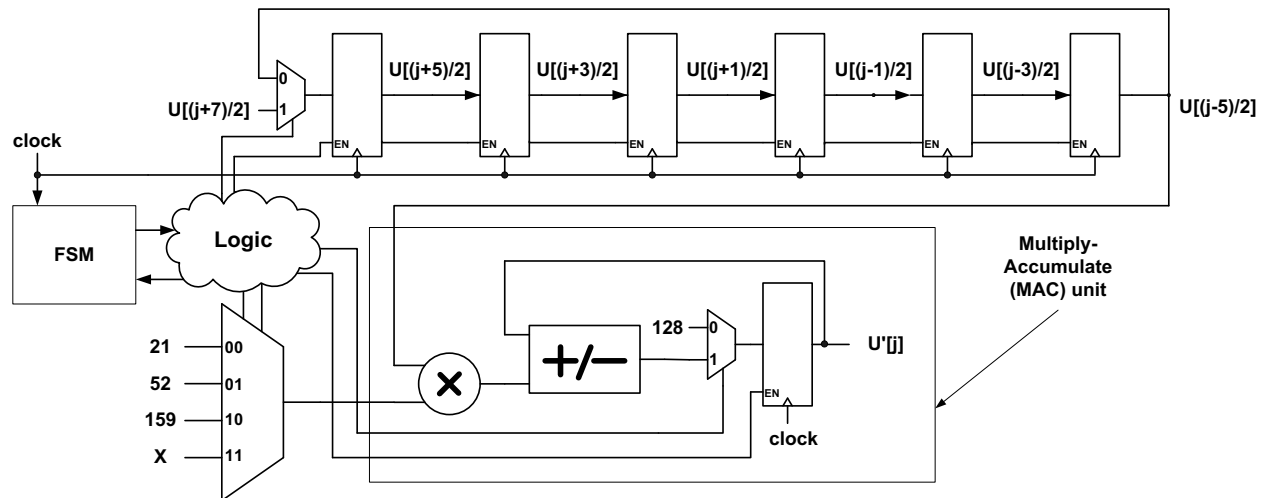


Figure 21. Resource efficient FIR filter implementation.

In addition to the FIR filter implementation, we are faced with the challenge of sharing a multiplier for colourspace conversion. Figure 22 shows the straightforward way to implement Equation 11, with the advantage of high throughput (one full RGB pixel per clock cycle). However, due to the constraints of using only **three multipliers** for both interpolation and colourspace conversion, this solution is not suitable for our implementation.

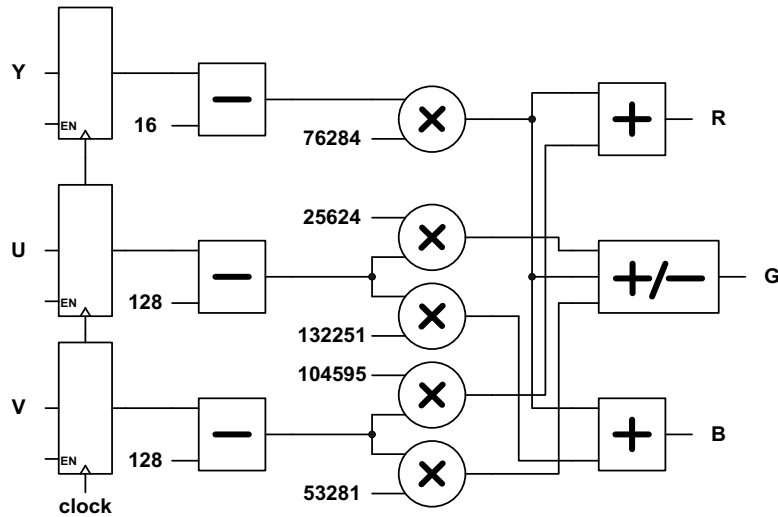


Figure 22. Resource inefficient colourspace conversion implementation.

In the same way that the MAC concept may be used to share multiplications for the FIR filters, it may be used for colourspace conversion. The data path for this is shown in Figure 23. In this case, using the mux between samples is sufficient since there are fewer samples from which to select, and the logic for the shift structure is not already in place, as was the case for the FIR filter.

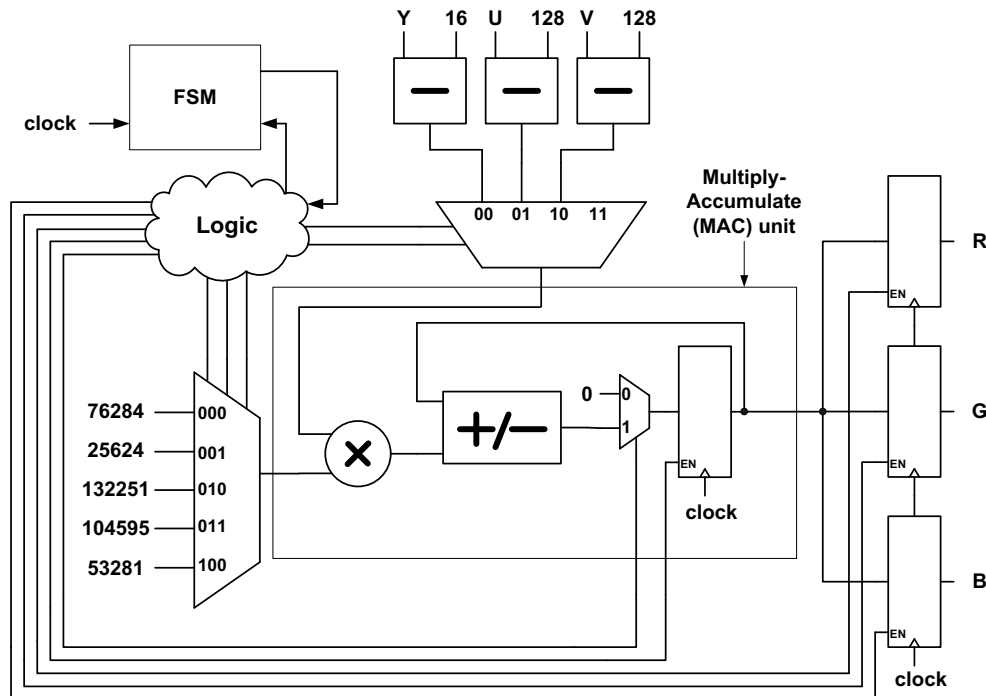


Figure 23. Resource efficient colourspace conversion implementation.

Table 2 provides a good starting point for a state table that can be used for milestone 1. You are strongly urged to use state tables throughout development, verification and debugging.

State code / clock cycle	0	1	2	3	4	5	...
SRAM_address							
SRAM_read_data							
SRAM_write_data							
SRAM_we_n							
R							
G							
B							
Y							
U'							
$U[(j+5)/2]$							
$U[(j-5)/2]$							
V'							
$V[(j+5)/2]$							
$V[(j-5)/2]$							
...							

Table 2. A starting point for the state table in milestone 1.

The average utilization of the 3 multipliers for milestone 1 must be at least 75%. The utilization of a multiplier is defined as the number of clock cycles a multiplier is used to perform calculations divided by the total number of clock cycles for the respective milestone. For clarification purposes, consider the following example on a hypothetical milestone whose duration is 100 clock cycles. Let's assume we use two multipliers. The first multiplier performs calculations for 70 clock cycles hence its utilization is 70%. The second multiplier performs calculations for 90 clock cycles hence its utilization is 90%. For this *hypothetical* milestone, the average utilization is 80%.

In addition to the above-mentioned utilization constraint, exploiting **the symmetry in the interpolation filter is mandatory**. The symmetry of the filter's coefficients from Equation 10 can be exploited as:

$$21 \left(U \left[\frac{j-5}{2} \right] + U \left[\frac{j+5}{2} \right] \right) - 52 \left(U \left[\frac{j-3}{2} \right] + U \left[\frac{j+3}{2} \right] \right) + 159 \left(U \left[\frac{j-1}{2} \right] + U \left[\frac{j+1}{2} \right] \right) \quad (13)$$

Note the 256 scale factor and the 128 offset from Equation (10) are not shown in the equation above, whose purpose is **only** to illustrate how symmetry can be exploited. Naturally, the 256 scale factor and the 128 offset from Equation (10) must be used when interpolating the odd samples for the U and V. Based on the above, it can be concluded that the interpolation for each of the odd samples from the U and V planes can be done with **only 3, rather than 6, multiplications**.

Note also that interpolated U' and V' can take values above 255 and below 0. These interpolated values should **not** be clipped before passing them to the colourspace converter (only the final R, G, and B values are clipped to 8 bits in the unsigned format in milestone 1).

Finally, as is the case for any digital design project, hardware resources should not be wasted. In your project report, you must justify your design decisions, e.g., the need for registers and their purpose.

4.3 Milestone 2: IDCT

Figure 24 shows the overview for milestone 2 which is the implementation of IDCT. As before, the shaded portions of the figure indicate what needs to be built or modified. In particular, notice that the FSM and the logic in the top state machine (bottom portion of the figure) do not need to be modified since, in the previous milestone, you have made provision for the decoder circuitry (all of which will eventually reside in the block currently labelled “Partial Decoder”) to access the SRAM.

As discussed in Section 4.1, the software model can generate the necessary initialization and verification data for the SRAM. The testbench uses this data for validation, and it can be sent to the board with a working design to get visual confirmation that the IDCT, upsampling, and colourspace conversion are working together.

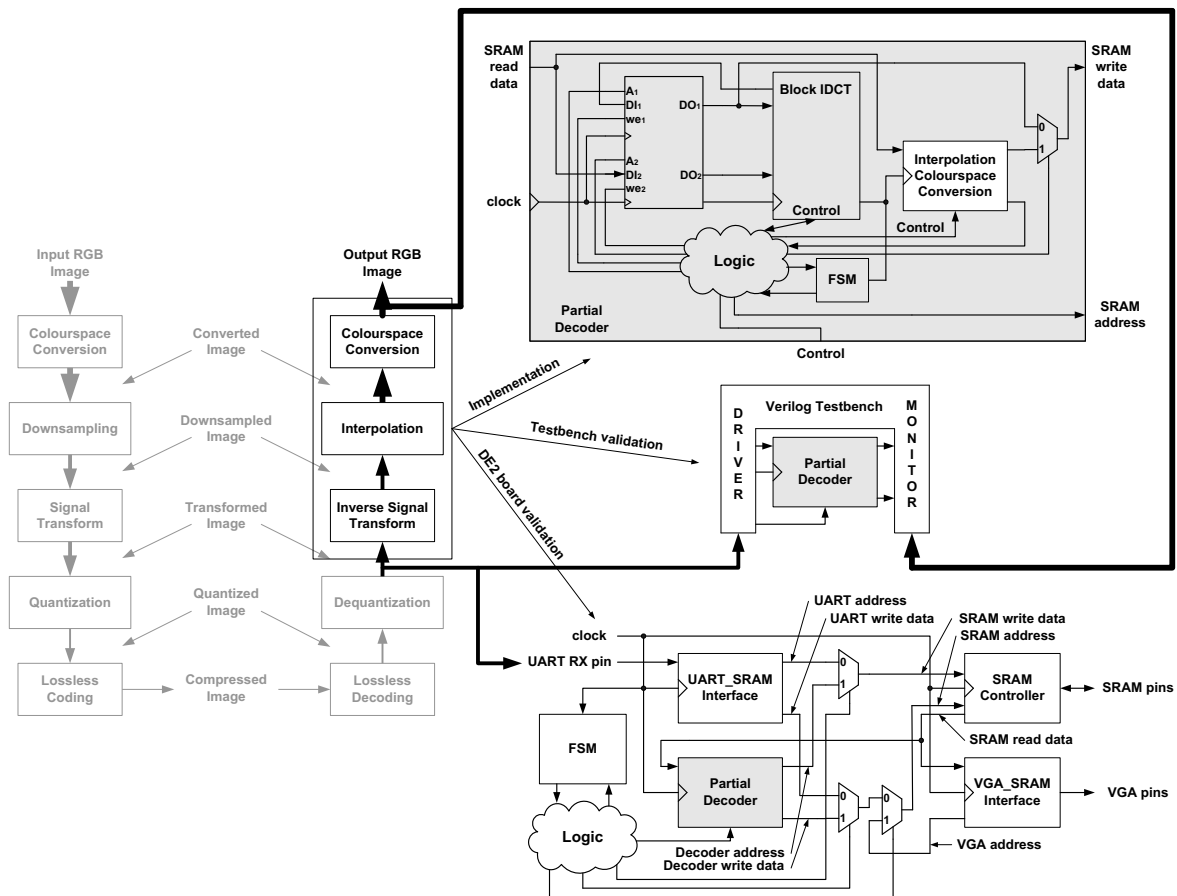


Figure 24. Overview of the IDCT implementation (milestone 2).

In Figure 25, the memory map for milestone 2 is provided, similar to that for milestone 1. The pre-IDCT data is stored as 1 value per memory location, whereas the post-IDCT data (YUV) is stored as 2 values per memory location. After IDCT is complete, the RGB data computed in milestone 1 will overwrite some pre-IDCT data (used to initialize the SRAM for milestone 2).

Like milestone 1, the blocks are divided into segments of Y, segments of (downsampled) U, and segments of (downsampled) V. However, the way the data is accessed is different. In milestone 2, we

want to fetch an entire 8x8 block of pre-IDCT data, perform IDCT and write it back to the YUV memory space. To determine the addressing for writing post-IDCT samples back to the memory, observe Figure 12 and note that the block from block row 0, block column 0 of the Y plane (Y 0,0) contains data from row 0, row1, ... row 7 of the image, all in columns 0 to 7 of the image. Extending this to the case of a 320x240 image, and noting that we have two samples per address, our write addressing for block 0 should be: 0 ... 3, ..., 1120 ... 1123 (as shown in Figure 26). For the second block of the first block row (Y 0,1), we have: 4 ... 7, ..., 1124 ... 1127. Once all the Y blocks have been processed, we proceed to U and V, with the important distinction that these two planes are downsampled and thus contain half the number of columns which Y contained. Thus, U 0,0 would reside in addresses (38400+): 0 ... 3, ..., 560 ... 563. The last block V 29,19 would reside in addresses (57600+): 18636 .. 18639, ..., 19196 ... 19199.

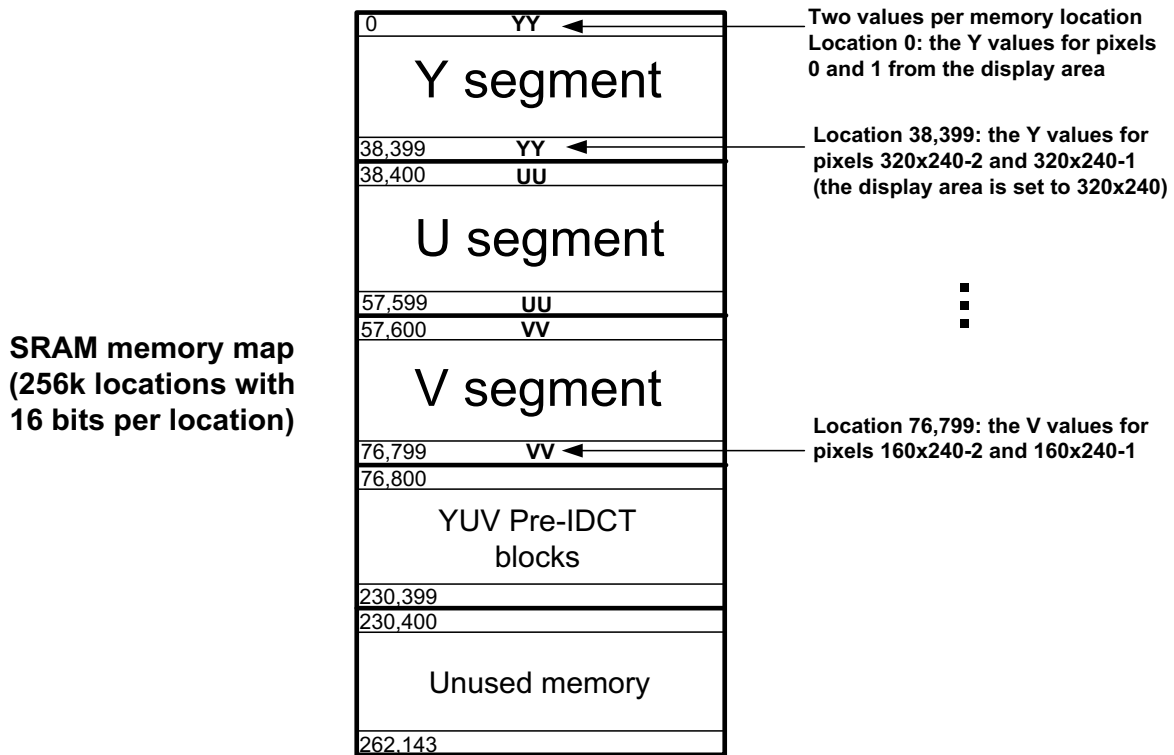


Figure 25. Memory layout for milestone 2.

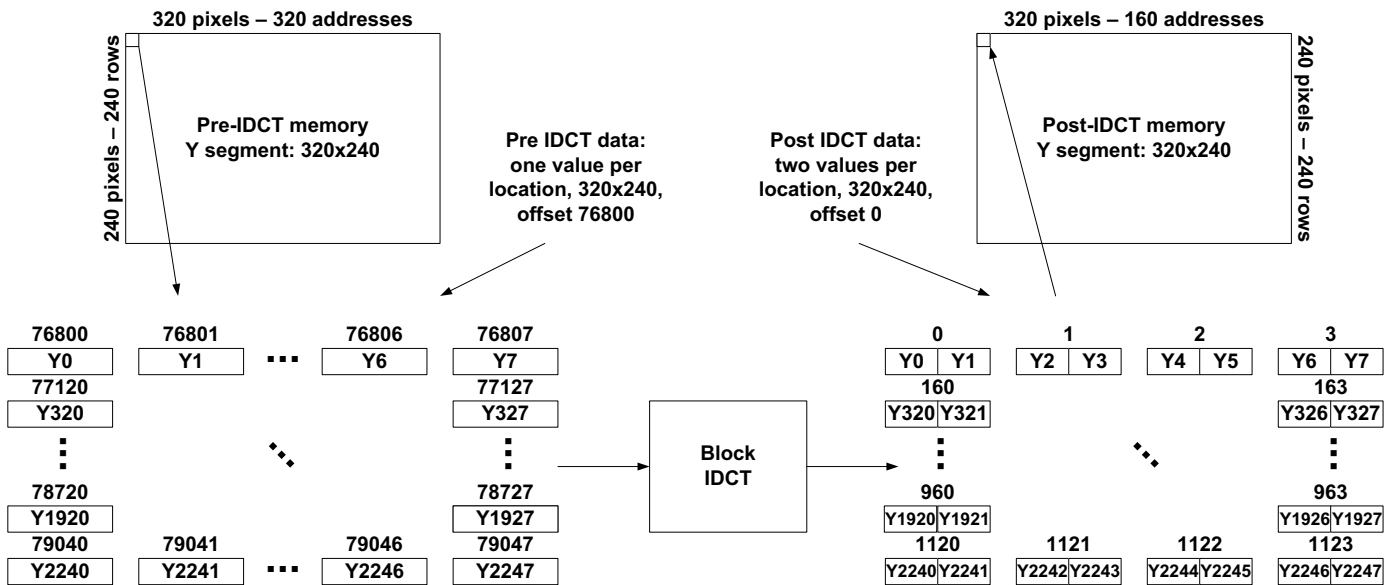


Figure 26. Reading pre-IDCT values and writing back YUV data to the external SRAM.

Like writing, the read addressing will also jump as we move from one pixel row to another within the same 8x8 block. However, now only 1 value is stored per memory location. To go across a row, the address simply increments. Within the same pixel column, to go down a pixel row (still within the same 8x8 block), the address increases by 320 (for Y, but only by 160 for U and V). Finally, we add the appropriate offset depending on our colour segment. Thus to read block Y 0,0 we use addresses (76800+): 0 ... 7, ..., 2240 ... 2247 (as in Figure 26). Likewise to read block Y 0,1 we use (76800+): 8 ... 15, ..., 2248 ... 2255. On the next block row, to read block Y 1,0 we use (76800+): 2560 ... 2567, ..., 4800 ... 4807. For the block U 0,0 we use (153600+): 0 ... 7, ..., 1120 ... 1127 (please note for U and V we have half the columns, hence there are half the number of blocks per row).

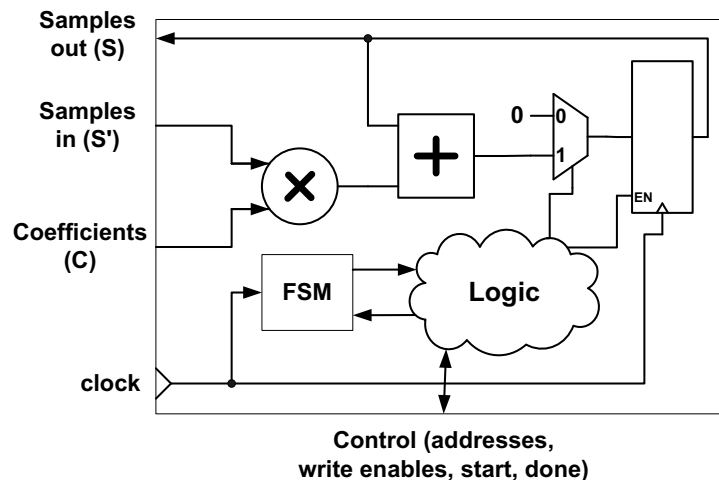


Figure 27. MAC unit for block IDCT.

Figure 27 shows a MAC unit, which may be used to calculate part of the IDCT on a single block of samples with a single multiplier. Recall that the IDCT is simply a pair of matrix multiplications. Since a

matrix multiplication with a constant coefficient matrix is nothing more than the sum of a number of products of the samples with constants, the MAC unit will meet the architectural need. The role of the FSM is to initialize the accumulator at the appropriate times and to set up the addresses properly to obtain the correct sample and coefficient in each clock cycle. Using such a unit, the pre-and post-matrix multiplication can be done concurrently (the number of multipliers for this milestone is clarified below).

To facilitate a simple hardware implementation, we can read from dual-port memories both the fixed-point cosine coefficients (Equation 9) and the data value (pre-IDCT values or intermediate values after the first matrix multiplication). It is your task to determine a suitable memory layout for the dual-port RAMs. Because you are given this freedom, you must provide the appropriate .hex or .mif file to specify the dual-port memory initialization, particularly for the fixed-point cosine coefficients.

To maximize the utilization of the embedded memories, it is recommended that both ports are used and that the memory is organized on 32 bits per location (the maximum width allowed). It is your choice whether the results after the second matrix multiplication are written directly to the SRAM or buffered first in a dual-port RAM.

In this milestone, ***you may use 3 dual-port RAMs (of 4Kb capacity and not more than 32 bits per location), and only 3 multipliers are allowed (with input operands on 32 bits or less).*** Note also that ***the average utilization of the 3 multipliers for milestone 2 must be at least 85%.***

You are encouraged to use a state table such as the one in Table 3 during your development for each matrix multiplication for the first few and the last few clock cycles.

Clock cycle	0	1	2	3	4	5	...
SRAM_address							
SRAM_read_data							
SRAM_write_data							
SRAM_we_n							
Address_0							
Data_in_0							
Write_en_0							
Data_out_0							
Address_1							
...							
S							
S'							
C							
...							

Table 3. A starting point for the state table in milestone 2.

4.4 Milestone 3: Lossless Decoding and Dequantization

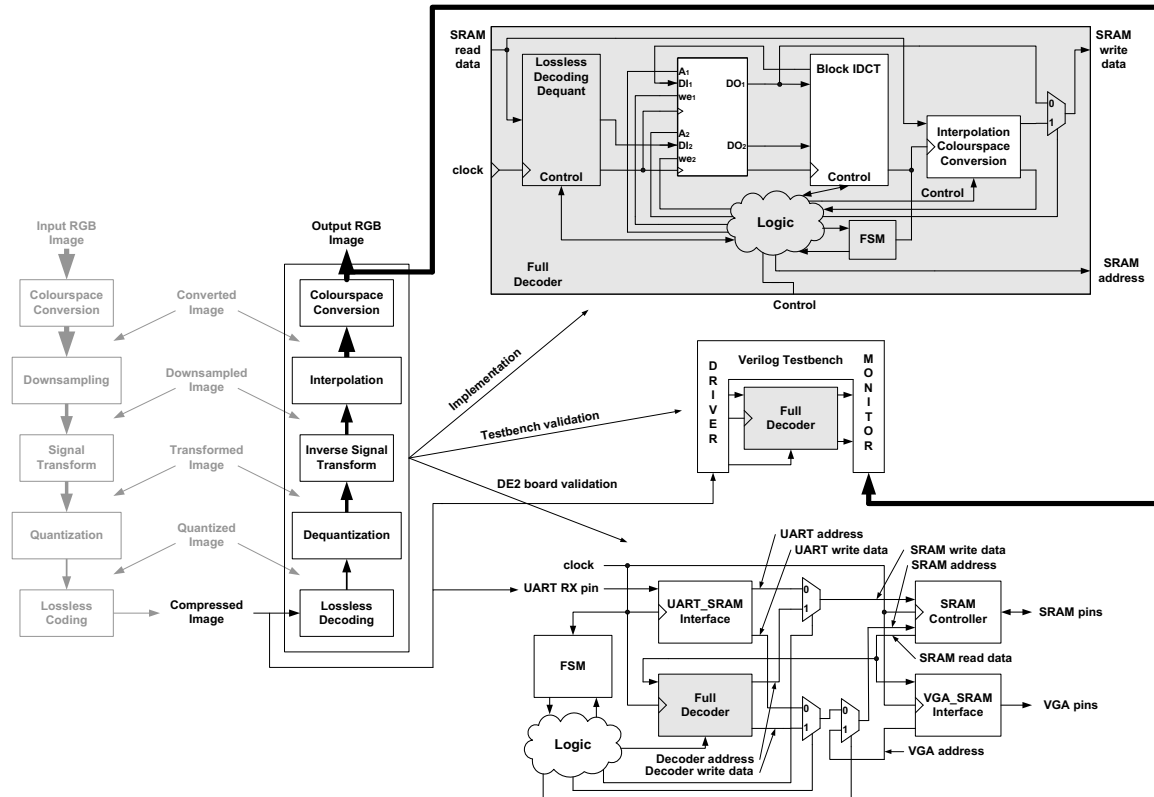


Figure 28. Lossless decoding, dequantization, and complete integration (milestone 3).

Finally, once colourspace conversion, interpolation and IDCT are complete, we can move on to lossless decoding, where the compressed bitstream itself is parsed, and the raw blocks are extracted and dequantized before being passed to the block IDCT, which completes the IDCT and writes the complete block of YUV data (belonging to the downsampled image of Figure 1) back to the SRAM. Once this milestone is complete, you will have a working image decoder.

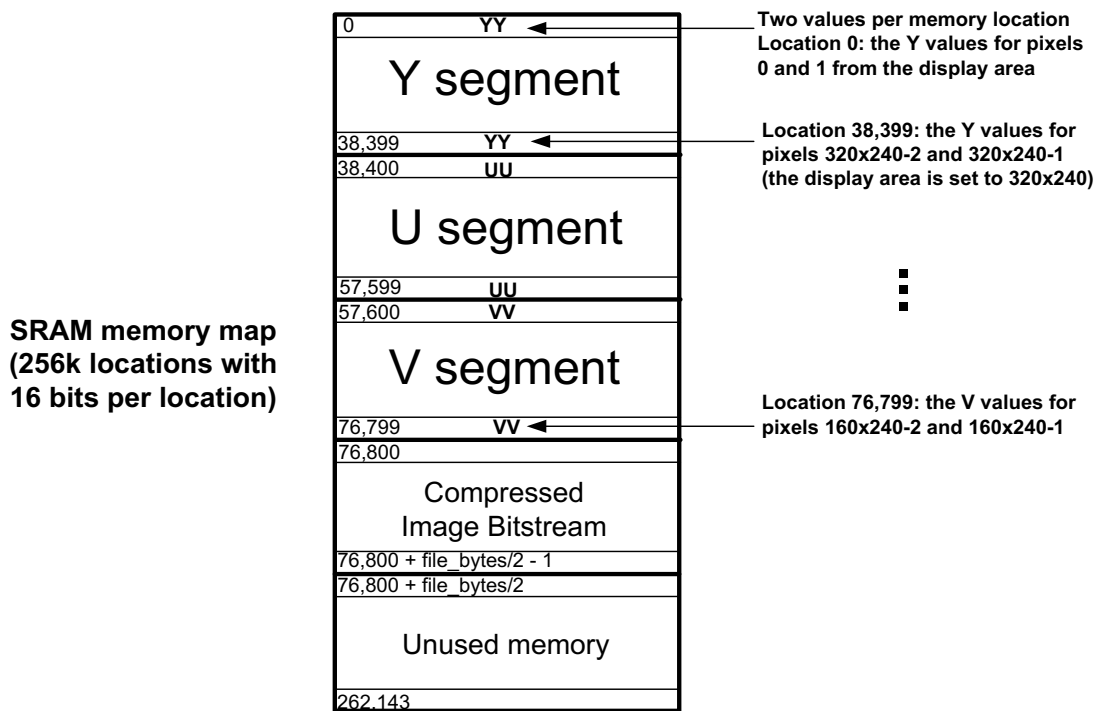


Figure 29. Memory layout for milestone 3.

Because the decoding process works on blocks of 8x8, after a block is decoded, it is immediately passed to block IDCT (milestone 2) by buffering into an embedded memory written to by milestone 3 and read from by milestone 2. The SRAM read address logic from milestone 2 will no longer be needed when milestones 2 and 3 work concurrently. Nevertheless, it is very important to note that you are strongly recommended to get milestone 3 working independently before integrating it to work concurrently with milestone 2. This will give you a chance to understand the functionality and work out the implementation details for the lossless decoder and the dequantizer, and only then can you focus on dealing with concurrency-related issues that will arise while integrating with milestone 2.

The memory layout for milestone 3 is shown above in Figure 29. No memory is allocated for the pre-IDCT data. Rather, this is processed by milestone 2 before being written back to the SRAM. After this, the RGB data generated by milestone 1 may overwrite the compressed image bitstream.

The basic decoding process is captured in Figure 8 at the beginning of the image decompression example. This flow chart can be used to construct a state table to guide your state machine implementation. Part of the challenge of lossless decoding is the implementation of a serializer, which enables the extraction of coefficient data by reading 16-bit words stored in the SRAM and delivering them to the decoder in a serial fashion. Figure 30 shows the basic architecture of the serializer. A custom FSM is required to drive the SRAM address and to increment it only if all the bits have been extracted from the current memory location. A 32-bit shift register can buffer the data on-chip, as codes and values may span two memory locations. In each clock cycle, the mux network will choose which bits from the location will be processed and shifted out of the 32-bit shift register. Note the shifting of bits in this 32-bit buffer may happen over more than one clock cycle, depending on your FSM (for example, we may shift out 6 bits by doing a shift by 3 on two consecutive clock cycles).

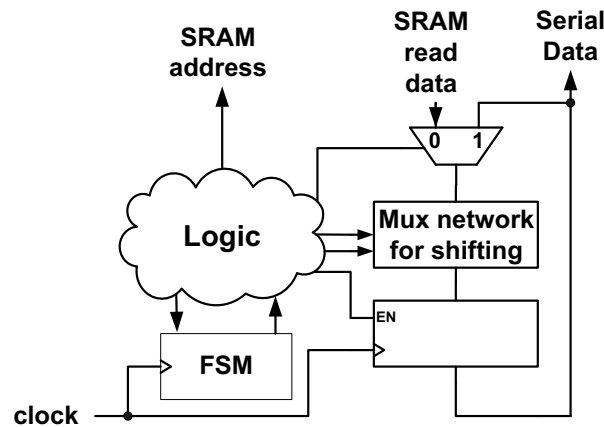


Figure 30. SRAM data is serialized for extracting coefficients.

Once the coefficients have been extracted, they must be dequantized before being placed in the dual-port RAM to await the block IDCT. Due to the particular structure of the two possible quantization matrices (as in Equation 3), which contain only powers of two, a barrel shifter structure can be used to implement the dequantization (the type of the quantization matrix must also be used as an input to this barrel shifter). Furthermore, note the scan sequence of Figure 4 must be implemented so that the values are delivered from the bitstream to the correct locations in the 8x8 block.

To facilitate easier integration between block IDCT and lossless decoding and quantization, ***you are allowed one extra dual-port RAM (of 4Kb capacity and not more than 32 bits per location). Thus a total of 4 dual-port RAMs can be used in milestones 2 and 3 together.*** This extra RAM can be used for loading the out-of-order pre-IDCT values, and then during the clock cycles when milestone 2 does not need to fetch pre-IDCT values, the content from this extra embedded RAM can be loaded into the embedded RAM shared between milestones 2 and 3. You may also write directly to the embedded RAM shared between milestones 2 and 3 (without the additional embedded RAM). However, this requires more design effort to ensure no scheduling conflicts between writing from milestone 3 and reading from milestone 2.

To handle a run of zeros, there are two basic approaches. One method is to write a sequence of zeros when it happens (in different memory locations and thus over several clocks). The other method is to initialize all values to zero, and then when a run of zeros is decoded, we can skip over that many memory locations.

Finally, successful implementation of lossless decoding and dequantization relies heavily on getting your state machine to work correctly and synchronizing the signals between it and the dual-port RAMs (used for IDCT), the serializer, the top-level state machine, etc.

ENJOY!

Appendix A: Summary of Data Flow and Computation

The 3DQ5 project must compile in Quartus and it is expected to operate correctly at the 50 MHz clock frequency.

Milestone 1

Input: data is represented as **8 bits unsigned**. Each SRAM location contains 2 values. There are separate sections for Y, U, and V. Within each section, data is stored in raster-scan order (across first, then down).

Computation: it is advised to do **all arithmetic on 32 bits signed**. Upsample the odd values of U and V (no clipping here). Each pixel is independently converted from YUV to RGB; after this, clipping to **8 bits unsigned** is done. In this milestone, you **must exploit the symmetry in the interpolation filter** and use only **3 multipliers** with at least 75% utilization.

$$U'[j] = \begin{cases} U\left[\frac{j}{2}\right] & j \text{ even} \\ (int)\frac{1}{256}\left(21U\left[\frac{j-5}{2}\right] - 52U\left[\frac{j-3}{2}\right] + 159U\left[\frac{j-1}{2}\right] + 159U\left[\frac{j+1}{2}\right] - 52U\left[\frac{j+3}{2}\right] + 21U\left[\frac{j+5}{2}\right] + 128\right) & j \text{ odd} \end{cases}$$

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = (int)\frac{1}{65536} \begin{bmatrix} 76284 & 0 & 104595 \\ 76284 & -25624 & -53281 \\ 76284 & 132251 & 0 \end{bmatrix} \left(\begin{bmatrix} Y \\ U' \\ V' \end{bmatrix} - \begin{bmatrix} 16 \\ 128 \\ 128 \end{bmatrix} \right)$$

Output: each of R, G and B are **8 bits unsigned**. Data for 2 pixels requires 3 SRAM locations to store. The memory layout of the RGB segment resembles lab 5 experiment 1 (see Figure 17 for starting address).

Milestone 2

Input: data is represented as **16 bits signed**.

Computation: it is advised to do **all arithmetic on 32 bits signed**. There are 2 matrix multiplications per output 8x8 block, after which each value is clipped to **8 bits unsigned**. Take note of the SRAM access pattern since data is processed in blocks of 8x8. **Up to 3 dual-port memories** (of 4Kb capacity and not more than 32 bits per location) can be used and it is your choice on how these are organized. In this milestone, you must use **3 multipliers** with at least 85% utilization.

$$S = (int)\frac{1}{65536} \left(C^T \times (int)\frac{1}{256} (S' C) \right)$$

Output: **8 bits unsigned** values are written to the SRAM, 2 values per SRAM location.

Milestone 3

Input: the .mic18 file is a **bitstream**; to get the next bit, we go from the most significant bit to the least significant bit within the same SRAM location, then only go to the following SRAM location.

Computation: decode each 8x8 block independently. Each prefix code may produce several values. Values are processed in the zigzag order for each block. Once the position in the 8x8 block is known for each value, dequantization is applied. The quantization matrix to be used must be extracted from the .mic18 header. **One extra dual-port memory** (same as the 3 memories from milestone 2) may be used.

Output: each dequantized value is represented as **16 bits signed**. Before you integrate milestone 3 into milestone 2, it is strongly recommended that you implement it separately, in which case, take note of the SRAM access pattern since data is processed in blocks of 8x8.